

Lecture 8: Optimization and hill-climbing

Artificial Intelligence

CS-GY-6613

Julian Togelius

julian.togelius@nyu.edu

On the menu:

- Optimization versus tree search
- Hill-climbing
- Simulated annealing
- Evolutionary algorithms

Tree search versus optimization

Local search algorithms

- In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use local search algorithms: keep a single "current" state, try to improve it

n-queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal




Hill-climbing

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a (highest-valued) successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```

Hill-climbing

- Can get stuck in local maxima/minima




Hill-climbing 8-queens

- h = number of pairs of queens that are attacking each other, either directly or indirectly. Here, $h=17$.

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

Local minimum



Hill-climbing variations

- Simple hill climber: generates one successor, accepts it if better than the current solution
- “Permissive” hill climber: accepts the successor, accepts it if *equal to or* better than the current solution
- Steepest ascent hill climber: generates *all* successors, chooses the best one

Optimization in general

- Optimal: no
- Complete: no
- Space: reasonable
- Time: who knows

Simulated annealing

- Do bad moves with decreasing probability

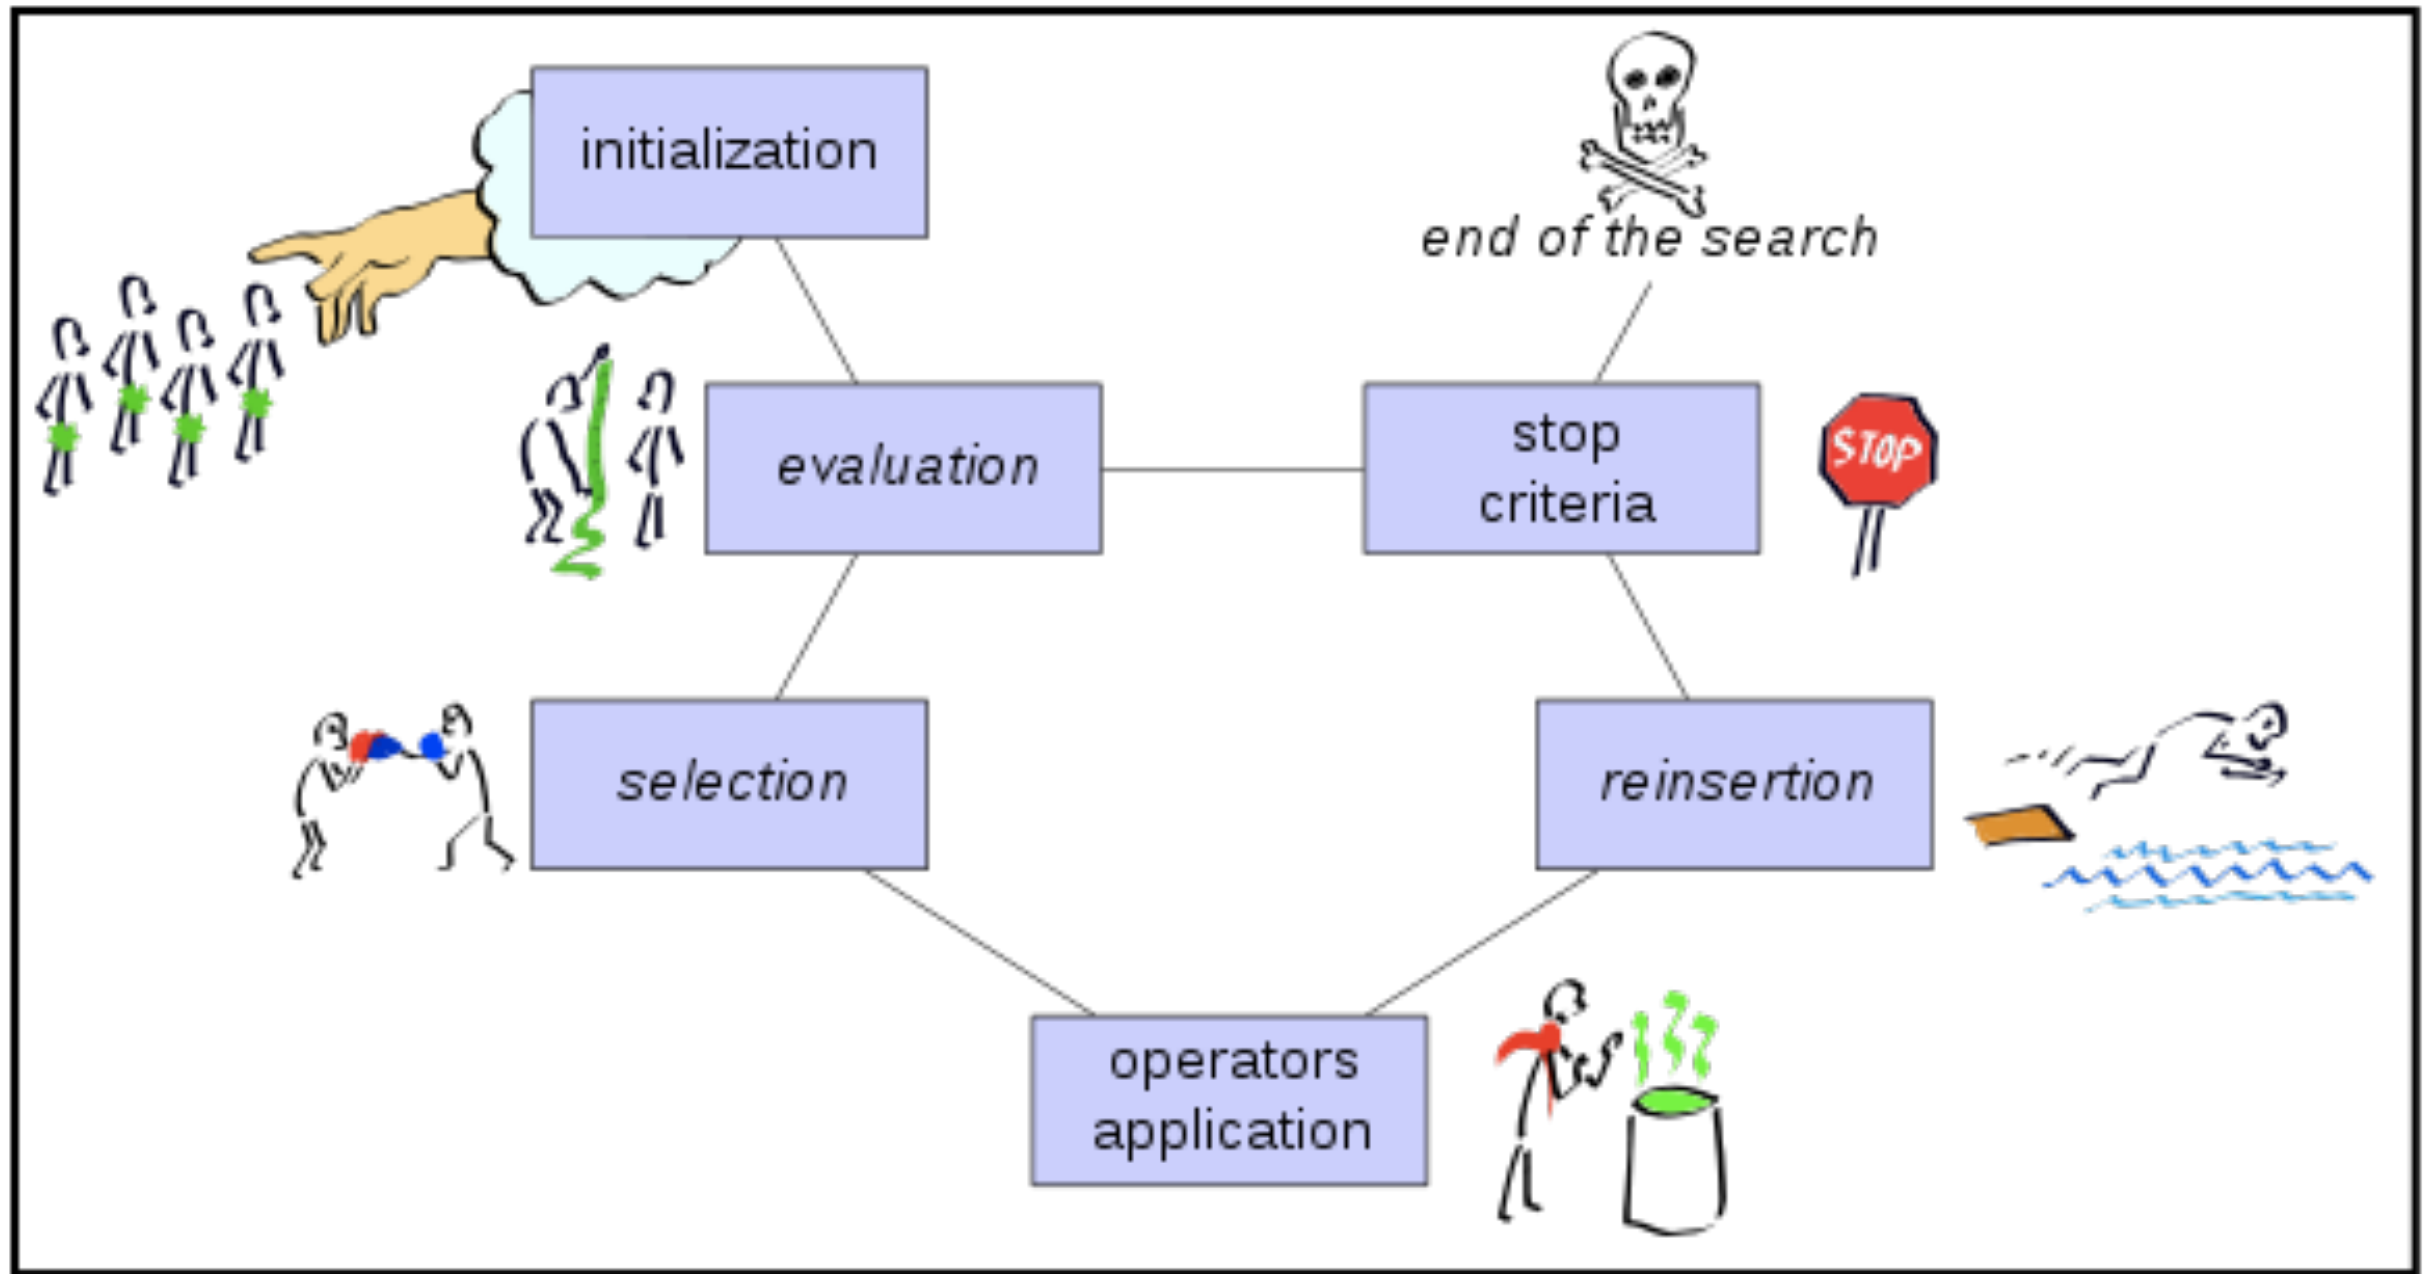
```
function SIMULATED-ANNEALING( problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to "temperature"
  local variables: current, a node
                    next, a node
                    T, a "temperature" controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

Simulated annealing

- One can prove: If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1

Evolutionary computation



General schema of an Evolutionary Algorithm (EA)

More about evolutionary
computation next time!