

# Lecture 9: Intro to Evolutionary Computation

Artificial Intelligence

CS-GY-6613

Julian Togelius

[julian.togelius@nyu.edu](mailto:julian.togelius@nyu.edu)

# Local search algorithms

- In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use local search algorithms: keep a single "current" state, try to improve it

# Hill-climbing

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```

# Types of hill-climbing

Hill-climbing is a somewhat generic concept, and there are several algorithms that fit the bill, such as:


- Stochastic hill-climbing: choose a random neighbor
- Steepest ascent hill-climbing: choose the best neighbor

# n-queens

- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal



# Local minimum



# Other ideas...

- Exhaustive search: try all configurations, one after another
- Random search: generate random configurations

# Random restarts

- When you stop making progress, start another hillclimber somewhere else
- Keep the best solution you found so far



# Simulated annealing

- Do bad moves with decreasing probability

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to "temperature"
  local variables: current, a node
                   next, a node
                   T, a "temperature" controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

# Limitations of local search

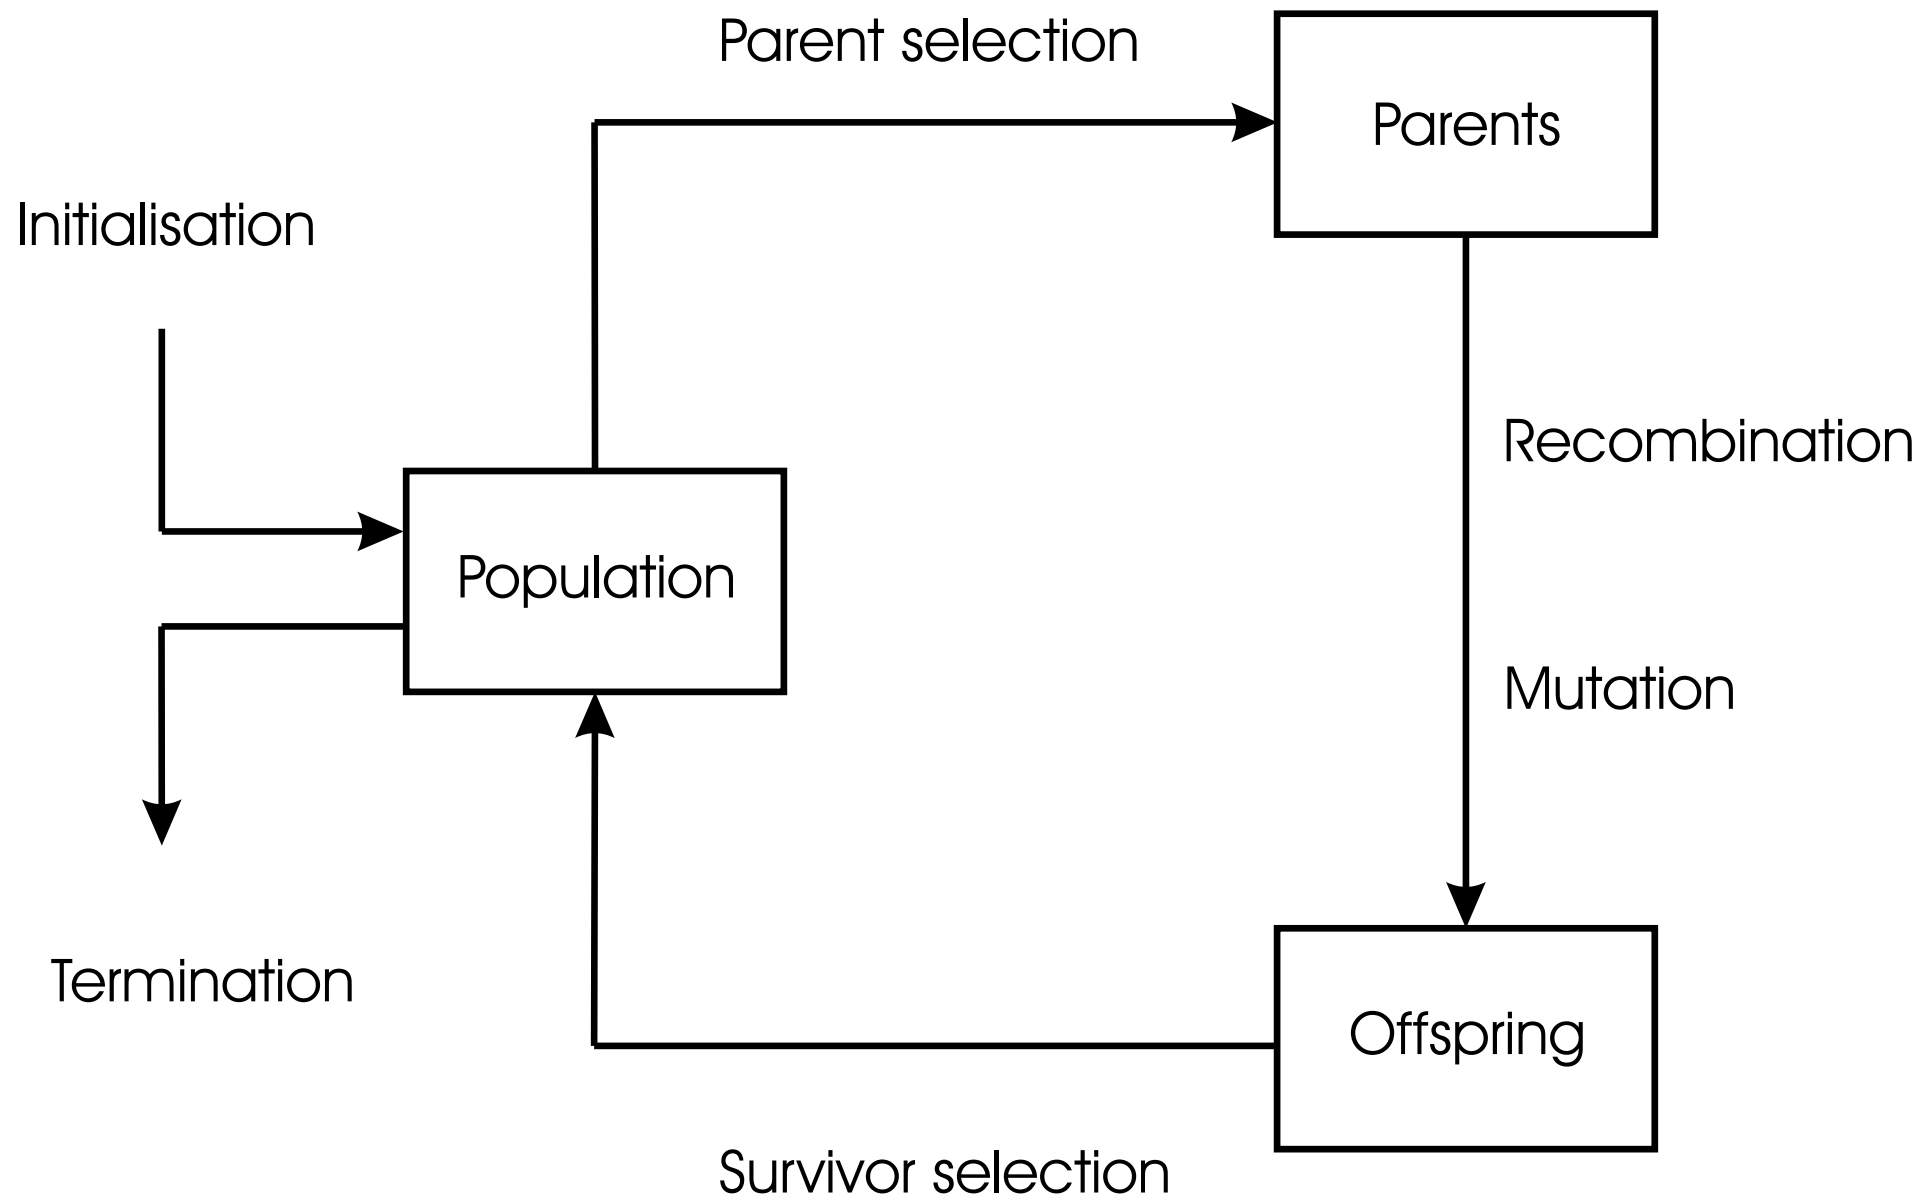
- Susceptible to local optima
- No sharing of information between parallel hill-climbing runs
- No possibility to learn from previous search experience

# Evolutionary algorithms

- Stochastic global optimisation algorithms
- Inspired by Darwinian natural evolution
- Extremely domain-general, widely used in practice



# Generic EA



# What's essential?

- Fitness evaluation
- Selection
  - Fitness influences how many offspring an individual has
- Variation
  - Could be mutation *and/or* crossover

# Generic EA

BEGIN

*INITIALISE* population with random candidate solutions;

*EVALUATE* each candidate;

REPEAT UNTIL ( *TERMINATION CONDITION* is satisfied ) DO

1 *SELECT* parents;

2 *RECOMBINE* pairs of parents;

3 *MUTATE* the resulting offspring;


4 *EVALUATE* new candidates;

5 *SELECT* individuals for the next generation;

OD

END

# The fitness landscape





# Simple $\mu+\lambda$ Evolution Strategy

- Create a population of  $\mu+\lambda$  individuals
- Each generation
  - Evaluate all individuals in the population
  - Sort by fitness
  - Remove the worst  $\lambda$  individuals
  - Replace with mutated copies of the  $\mu$  best

# Simple $\mu+\lambda$ ES (concrete)

- Create a population of  $\mu+\lambda$  individuals (e.g. 50+50) represented as e.g. vectors of real numbers in  $[0..1]$
- Each generation (until 100 generations)
  - Evaluate all individuals in the population
  - Sort by fitness (e.g. win rate or score for the game; higher is better)
  - Remove the worst  $\lambda$  individuals
  - Replace with mutated copies of the  $\mu$  best (mutate through Gaussian mutation with mean 0 and s.d. 0.1)

Stochastic hillclimber =  
evolution strategy with  $\mu=\lambda=1$

# Example: OneMax

- Representation: binary strings of length 5
- Examples: 00000, 01001, 00100, 11101, 11111
- Fitness function: count the number of ones
- Mutation: flip a random bit

# Example: OneMax

- Initial population: 00100, 11001, 10010, 01010
- Gen 1: 11001 : 3, 10010 : 2, 01010 : 2, 00100 : 1  
Selection: 11001 : 3, 10010 : 2
- Gen 2: 11101 : 4, 11001 : 3, 10010 : 2, 00010 : 1  
Selection: 11101 : 4, 11001 : 3
- Gen 3: 11111 : 5, 11101 : 4, 11101 : 4, 11001 : 3

# Generally, you need

- A solution representation
- Variation operators (mutation and/or crossover)
- A fitness (evaluation) function