

Lecture 12: Alpha-Beta and state evaluation

Artificial Intelligence

CS-GY-6613

Julian Togelius

julian.togelius@nyu.edu

Properties of Minimax


- *Complete?* Yes (if tree is finite)
- *Optimal?* Yes (against an optimal opponent)
- *Time complexity?* $O(b^m)$
- *Space complexity?* $O(bm)$ (depth-first exploration)

For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible


α - β pruning

- Can we improve this?
- Idea: don't consider branches of the tree that cannot lead to a better outcome than those that we have already explored


α - β pruning




α - β pruning




α - β pruning



α - β pruning



α - β pruning




Properties of α - β

- Pruning does **not** affect final result
- Good move ordering improves effectiveness of pruning
- With "perfect ordering", time complexity = $O(b^{m/2})$
 - doubles depth of search
- A simple example of the value of reasoning about which computations are relevant (metareasoning)

Why is it called α - β ?

- α is the value of the best (i.e. highest-value) choice found so far at any choice point along the path for *max*
- If v is worse than α , *max* will avoid it
> prune that branch
- Define β similarly for *min*



function ALPHA-BETA-SEARCH($state$) **returns** *an action*

inputs: $state$, current state in game

$v \leftarrow \text{MAX-VALUE}(state, -\infty, +\infty)$

return the *action* in SUCCESSORS($state$) with value v

function MAX-VALUE($state, \alpha, \beta$) **returns** *a utility value*

inputs: $state$, current state in game

α , the value of the best alternative for MAX along the path to $state$

β , the value of the best alternative for MIN along the path to $state$

if TERMINAL-TEST($state$) **then return** UTILITY($state$)

$v \leftarrow -\infty$

for a, s in SUCCESSORS($state$) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v

```

function MIN-VALUE( $state, \alpha, \beta$ ) returns a utility value
  inputs:  $state$ , current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to  $state$ 
            $\beta$ , the value of the best alternative for MIN along the path to  $state$ 

  if TERMINAL-TEST( $state$ ) then return UTILITY( $state$ )
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS( $state$ ) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 

```


Still, there is no time...

- Suppose we have 100 secs, and can explore 10^4 nodes/sec
→ 10^6 nodes per move
(a far cry from 35^{100} ...)
- Standard approach: use an *evaluation function* (heuristic)
 - Cut off search and treat states as end states
- Either at the same depth for all branches, or use a *cutoff test* such as in quiescence search

Evaluation functions

- Simplest: number of white pieces - number of black pieces
- More complex: assign values to piece types
- Even more complex: count number of threats, try to recognize known positions
- Generally: Linear weighted sum of features
- Neural network
- Read more: *Blondie24* by David Fogel

Piece weights?



- MinimaxCutoff is identical to MinimaxValue except
 - Terminal? is replaced by Cutoff?
 - Utility is replaced by Eval
- Does it work in practice?
 - $b^m = 10^6$, $b=35$ means $m=4$
- 4-ply lookahead is (in general) a hopeless chess player!
 - 4-ply \approx human novice
 - 8-ply \approx old-school Chess program, human master
 - 12-ply \approx Deep Blue, Kasparov

Some deterministic two-player games

- Chess: Kasparov vs Deep Blue 1997
- Checkers: Chinook defeated Tinsley (grand champion) 1994
 - Solved 2007
- Othello: computers vastly better than humans
- Go: AlphaGo defeated Lee Sedol 2016



Google DeepMind

Challenge Match

8 - 15 March 2016



AlphaGo



AlphaGo

- Go has enormous branching factor
- Very hard to come up with an accurate state evaluation function
- AlphaGo is a combination of:
 - Monte Carlo Tree Search (next lecture)
 - Supervised learning in neural networks (October)
 - Reinforcement learning through self-play (November)