

Lecture 13:

Monte Carlo Tree Search

Artificial Intelligence

CS-GY-6613

Julian Togelius

julian.togelius@nyu.edu

Let's Go!



Rules

- Players take turns to place stones at intersections
- Every stone remaining on the board must have at least one open "point" (an intersection, called a "liberty") directly next to it (up, down, left, or right)
- The stones on the board must never repeat a previous position of stones

Branching factors

- Chess: ~35
- Go: ~350
- Hex: ~100
- Arimaa: ~13000
- Pac-Man: 4
- Halo: ...infinite?
- Most games: lots

Heuristic evaluation functions

- Chess: count the number of white vs black pieces, who holds key positions etc.
- Go: ???
- Pac-Man: number of pills eaten, proximity to ghosts?
- Halo: ???


Go MCTS!

| Year | Program | Description | Elo |
|------|-------------|---|------|
| 2006 | INDIGO | Pattern database, Monte Carlo simulation | 1400 |
| 2006 | GNU GO | Pattern database, α - β search | 1800 |
| 2006 | MANY FACES | Pattern database, α - β search | 1800 |
| 2006 | NEUROGO | TDL, neural network | 1850 |
| 2007 | RLGO | TD search | 2100 |
| 2007 | MoGo | MCTS with RAVE | 2500 |
| 2007 | CRAZY STONE | MCTS with RAVE | 2500 |
| 2008 | FUEGO | MCTS with RAVE | 2700 |
| 2010 | MANY FACES | MCTS with RAVE | 2700 |
| 2010 | ZEN | MCTS with RAVE | 2700 |

What is MCTS?

- A way of selecting the next action
- A statistical tree-search method with rollouts rather than function evaluations
- Builds unbalanced trees

Bandit problems



- At each step, pull one arm
- Noisy/random reward signal in the range $[0..1]$
- Different average reward
- Task: maximise reward (minimise regret)

Which arm to pull?

Which arm to pull?

- Pull all arms equally often?
- Only pull the arm that has given the best results so far?
- Mostly pull the “best” arm, but sometimes one of the others?
- An example of the exploration/exploitation dilemma
- Principled solution?

Which arm to pull?

Flat Monte Carlo

Share trials uniformly
between arms

ϵ -Greedy

$P(1 - \epsilon)$ – Best arm so far
 $P(\epsilon)$ – Random arm

UCB1 (Auer et al (2002)).

Choose arm j so as to
maximise:

$$\bar{X}_j + \sqrt{\frac{2 \log n}{T_j(n)}}$$

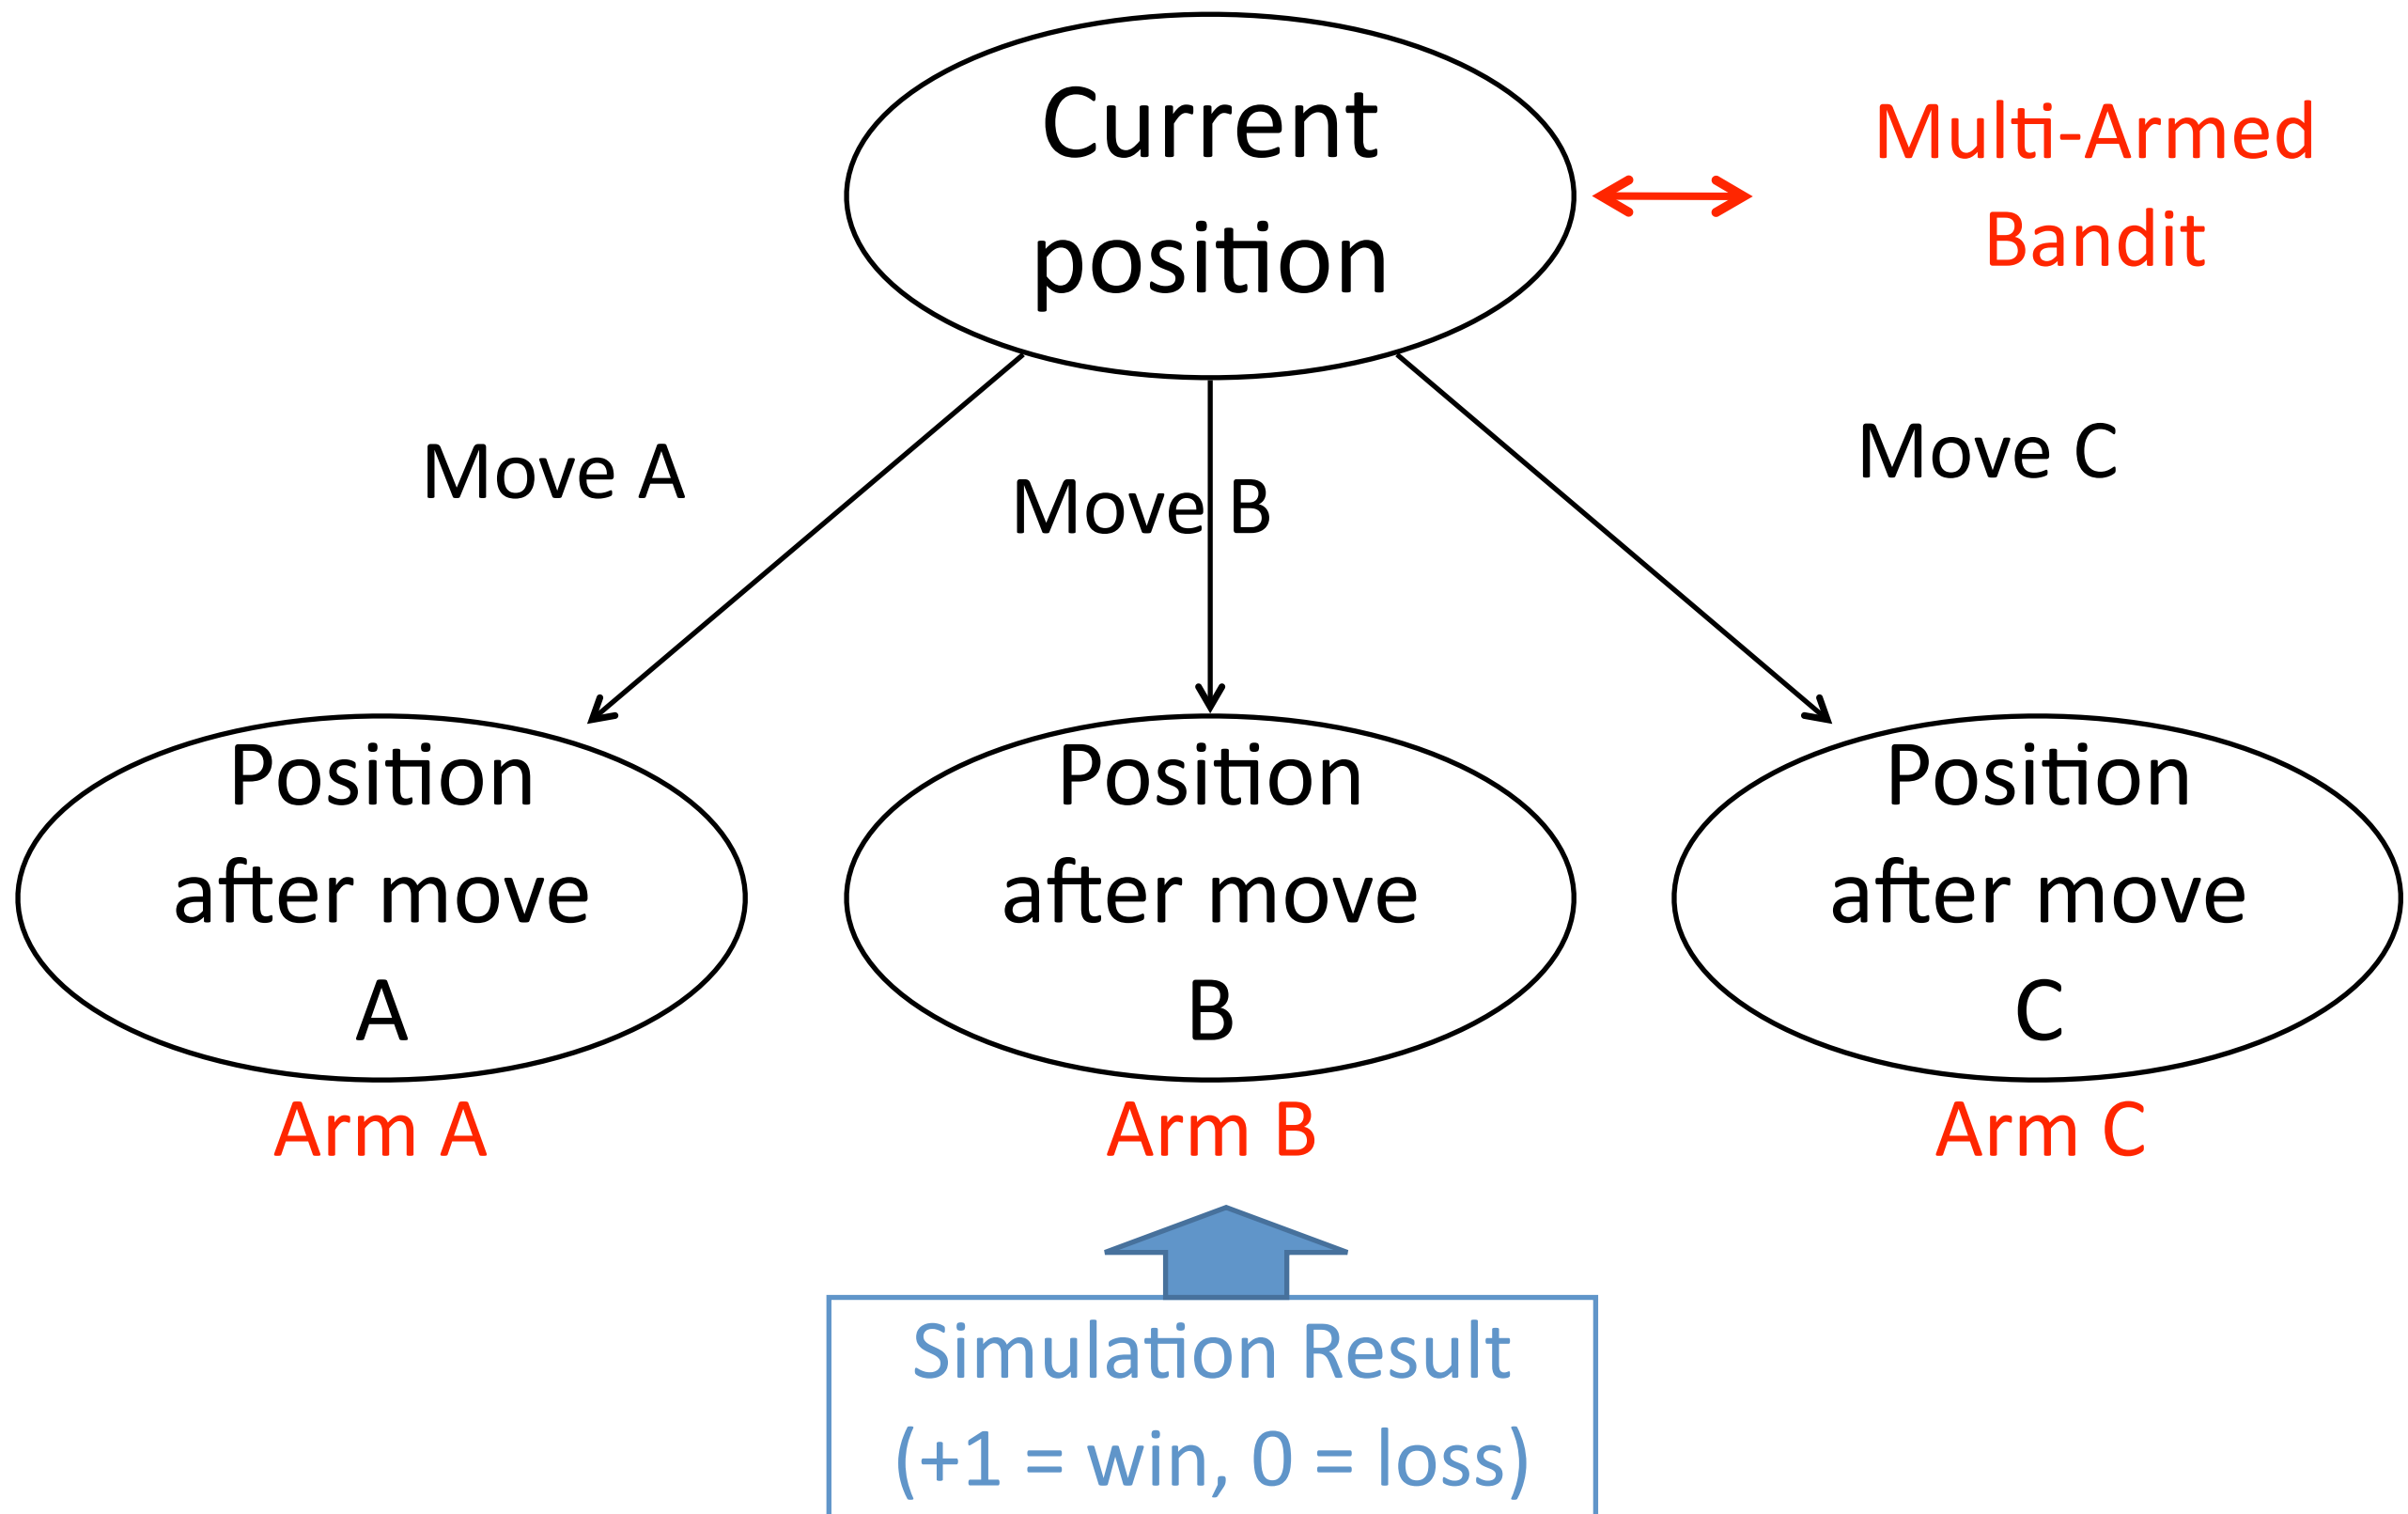
Mean
so far

Upper bound
on variance

n = number of plays so far

$T_j(n)$ = number of times arm j was pulled

Game Decisions



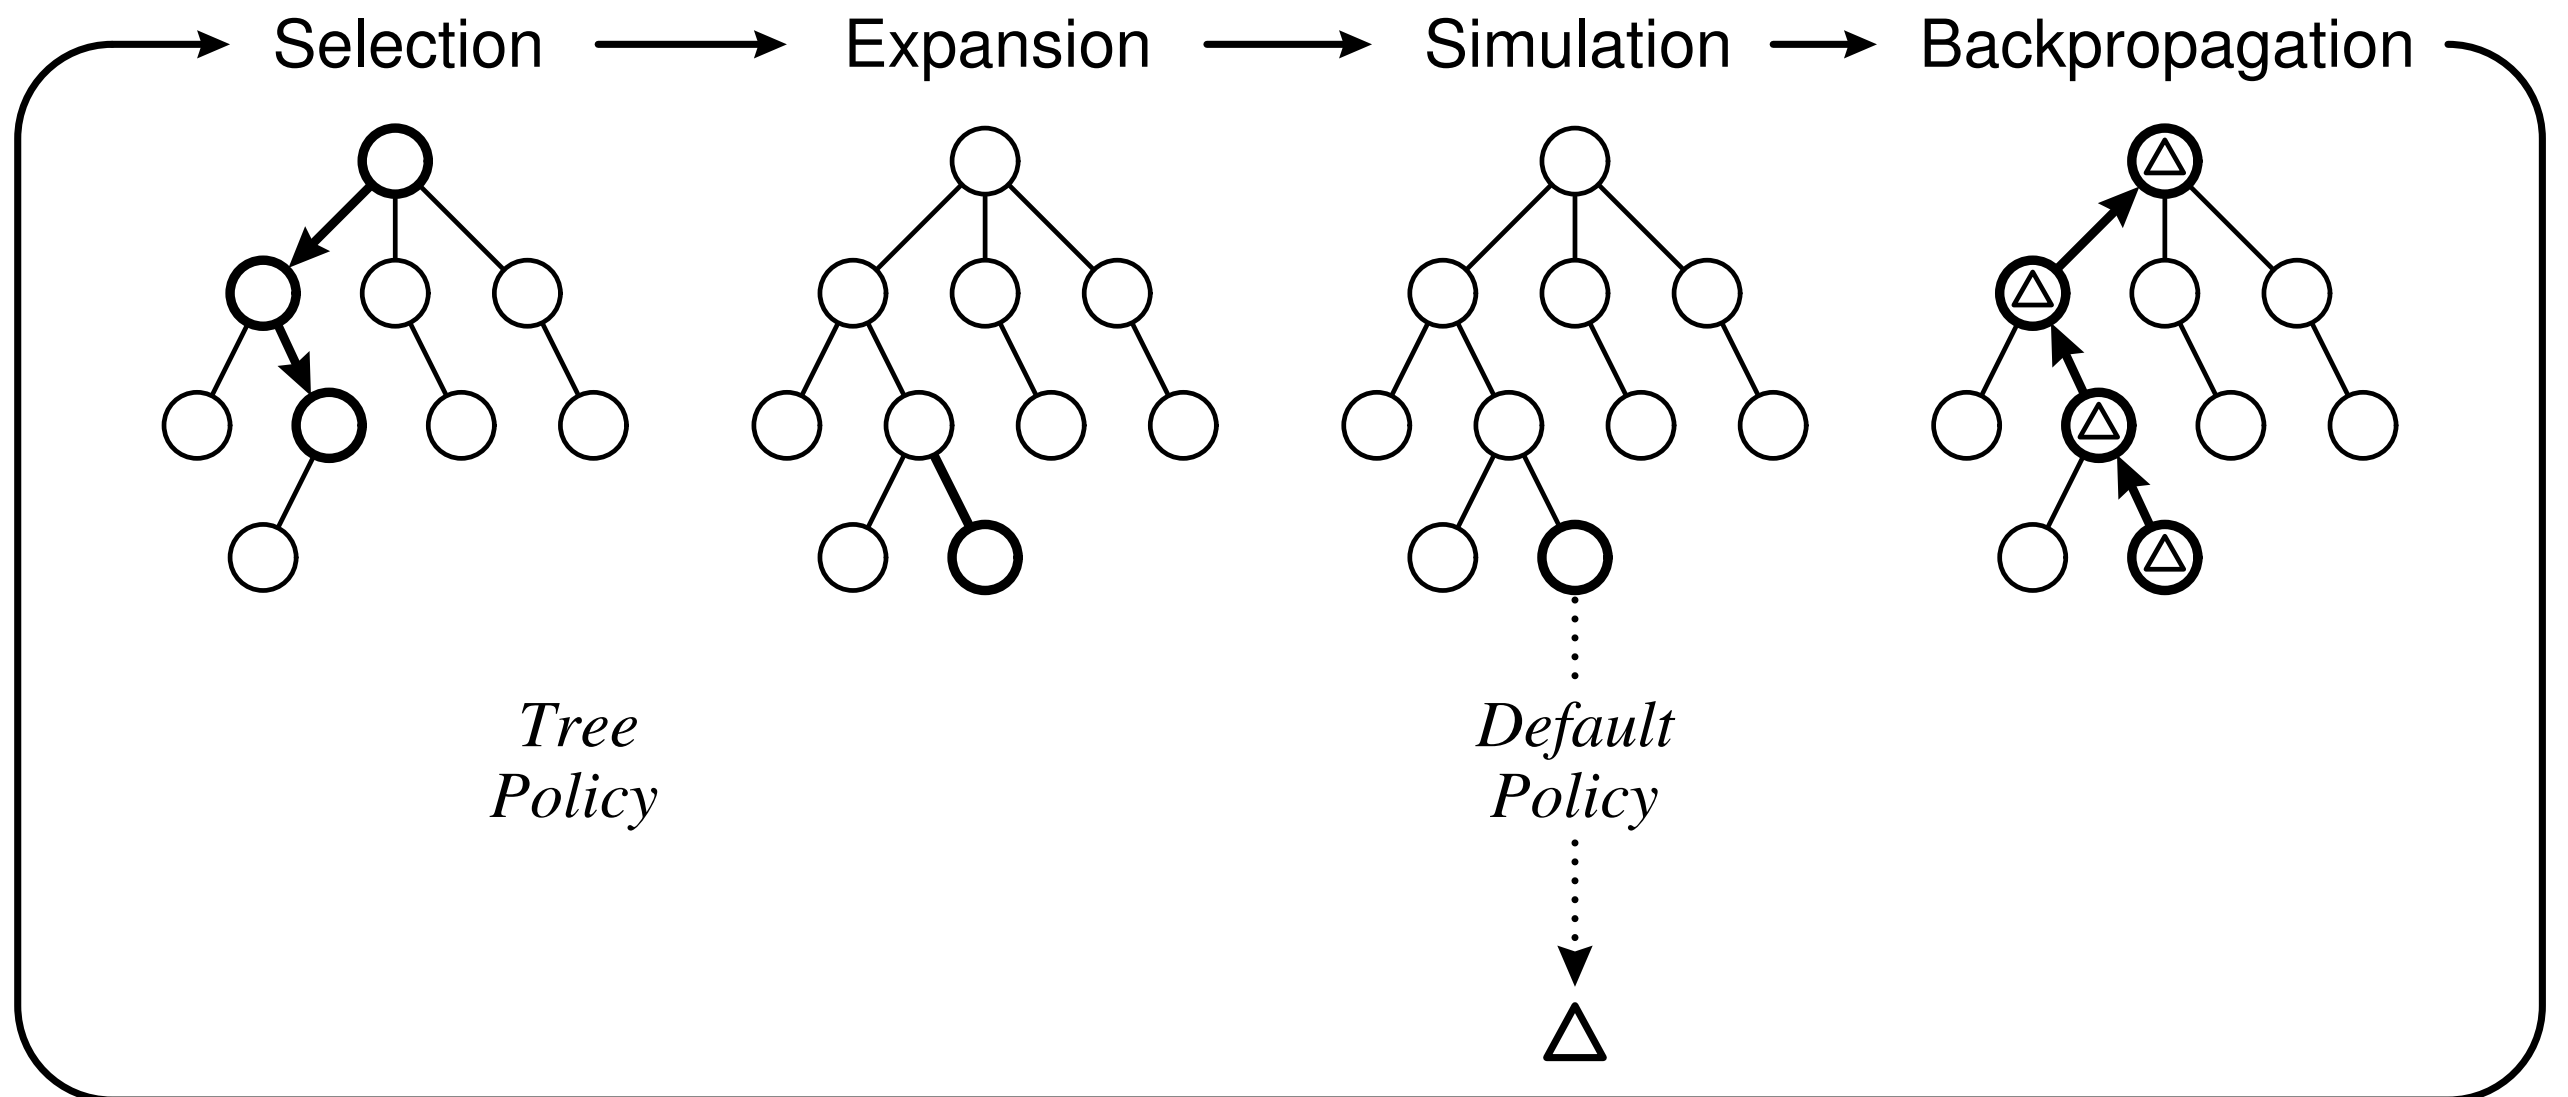
UCB

- Anytime – stop whenever you like
- UCB1 formula minimises regret
- Grows like $\log(n)$
- Needs only game rules:
 - Move generation
 - Terminal state evaluation
- Surprisingly effective, but...
... doesn't look ahead

UCT (UCB on trees)

- Anytime
- Scalable
- Tackle complex games better than before
- May be logarithmically better with increased CPU
- No need for heuristic function
- Though often better with one

MCTS general idea



MCTS algorithm

- Tree policy
 - Expand
 - Best Child (UCT Formula)
- Default Policy
- Back-propagate

Algorithm 1 General MCTS approach.

```
function MCTSSEARCH( $s_0$ )  
  create root node  $v_0$  with state  $s_0$   
  while within computational budget do  
     $v_l \leftarrow \text{TREEPOLICY}(v_0)$   
     $\Delta \leftarrow \text{DEFAULTPOLICY}(s(v_l))$   
    BACKUP( $v_l, \Delta$ )  
  return  $a(\text{BESTCHILD}(v_0))$ 
```

Tree policy

```
function TREEPOLICY( $v$ )  
  while  $v$  is nonterminal do  
    if  $v$  not fully expanded then  
      return EXPAND( $v$ )  
    else  
       $v \leftarrow$  BESTCHILD( $v, C_p$ )  
  return  $v$ 
```

Note that node selected for expansion does not need to be a leaf of the tree (the nonterminal test refers to the game state)

Tree expansion

```
function EXPAND( $v$ )  
  choose  $a \in$  untried actions from  $A(s(v))$   
  add a new child  $v'$  to  $v$   
    with  $s(v') = f(s(v), a)$   
    and  $a(v') = a$   
return  $v'$ 
```

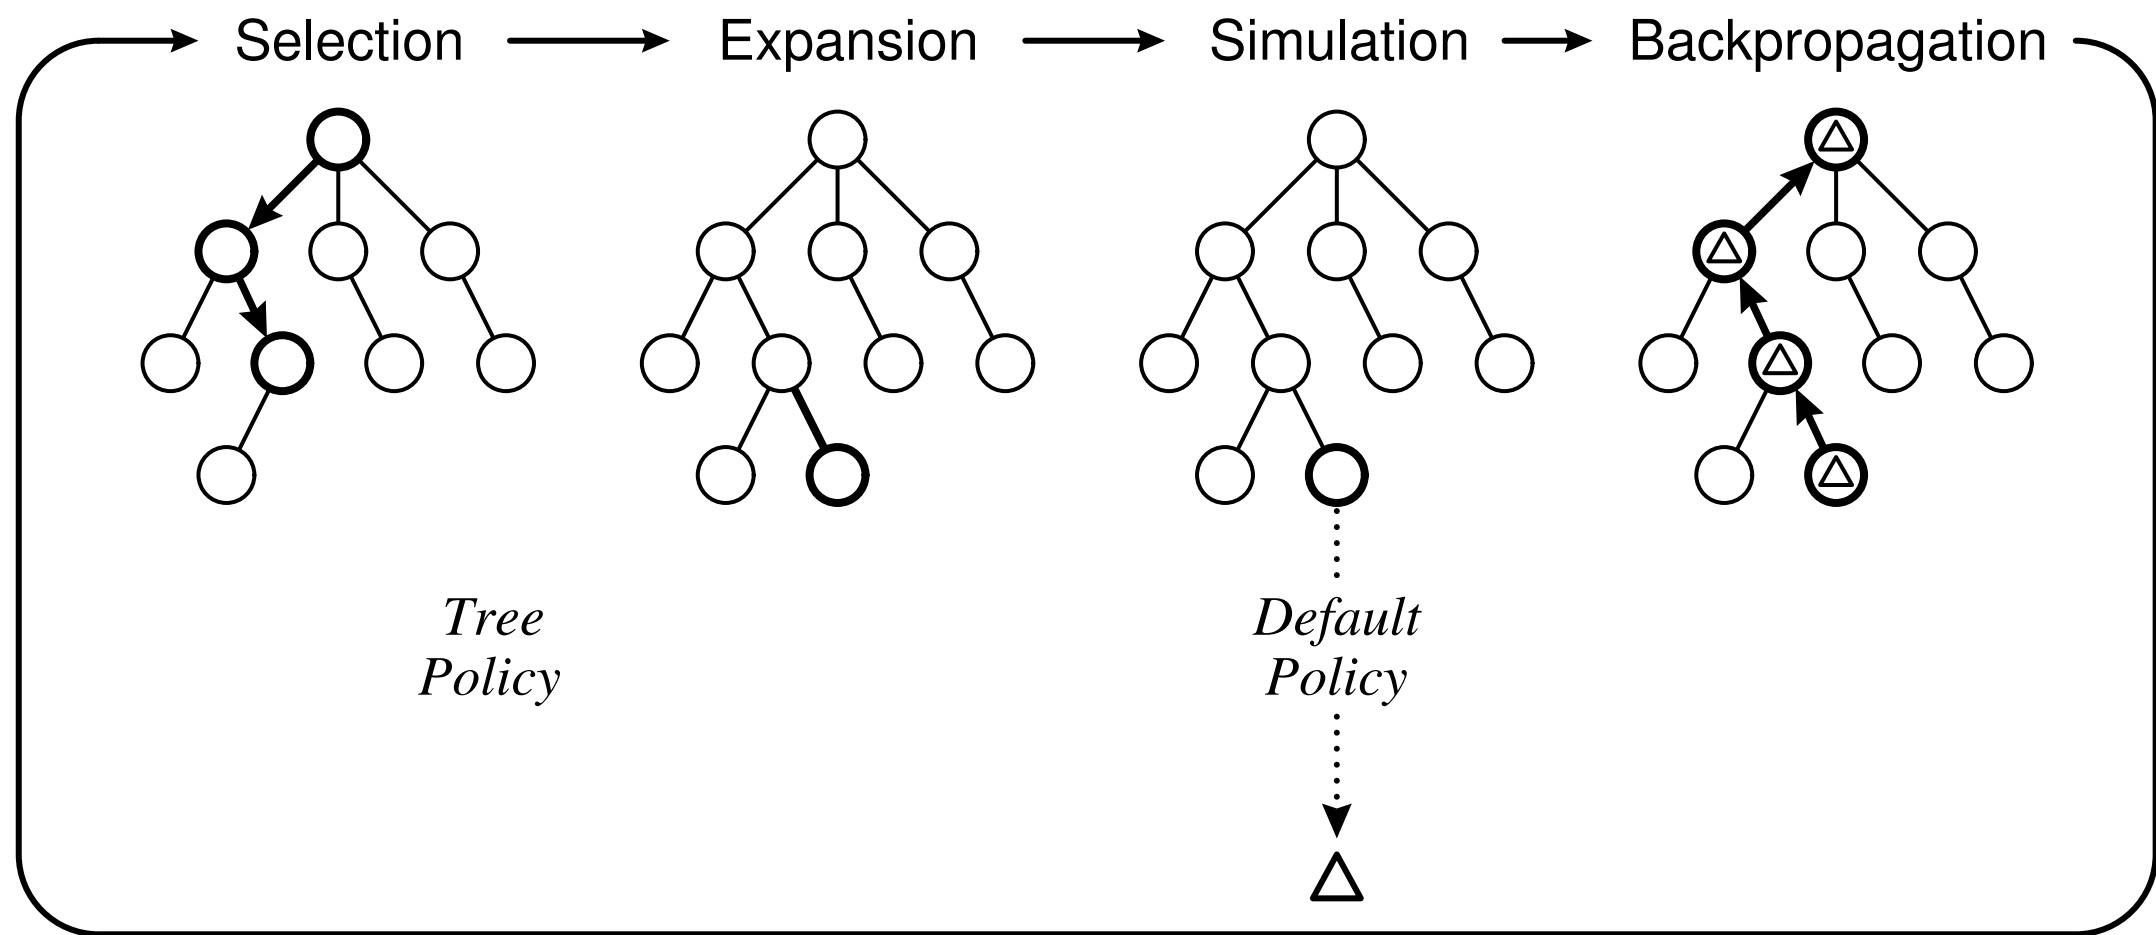
Best child (UCT)

function BESTCHILD(v, c)

return $\arg \max_{v' \in \text{children of } v} \frac{Q(v')}{N(v')} + c \sqrt{\frac{2 \ln N(v)}{N(v')}}}$

- Standard UCT equation (compare UCB)
- Higher values of c lead to more exploration

MCTS general idea



- Tree policy: choose which node to expand (not necessarily leaf of tree)
- Default (simulation) policy: random playout until end of game

Default policy (rollout)

```
function DEFAULTPOLICY( $s$ )  
  while  $s$  is non-terminal do  
    choose  $a \in A(s)$  uniformly at random  
     $s \leftarrow f(s, a)$   
  return reward for state  $s$ 
```


- Each time a node is added to the tree, the default policy plays out until the terminal state of the game
- The standard is to do this uniformly randomly

Backup

```
function BACKUP( $v, \Delta$ )  
  while  $v$  is not null do  
     $N(v) \leftarrow N(v) + 1$   
     $Q(v) \leftarrow Q(v) + \Delta(v, p)$   
     $v \leftarrow$  parent of  $v$ 
```

- v is the new node added to the tree by the tree policy
- Back up the values from the added node up the tree to the root

MCTS builds asymmetric trees



Enhancements

- Selection/expansion
 - All moves as first (AMAF) / RAVE
 - Bandit enhancements
 - Parameter tuning
- Simulation enhancements
 - Make the simulation more realistic (often costly and results may vary)

Non-determinism and incomplete information

- Many games are not deterministic
- Many games are only partially observable
- Determinization: create separate nodes for each random outcome
 - Huge branching factor
- Cheat?