

# Lecture 19:

# Multilayer Perceptrons


Artificial Intelligence

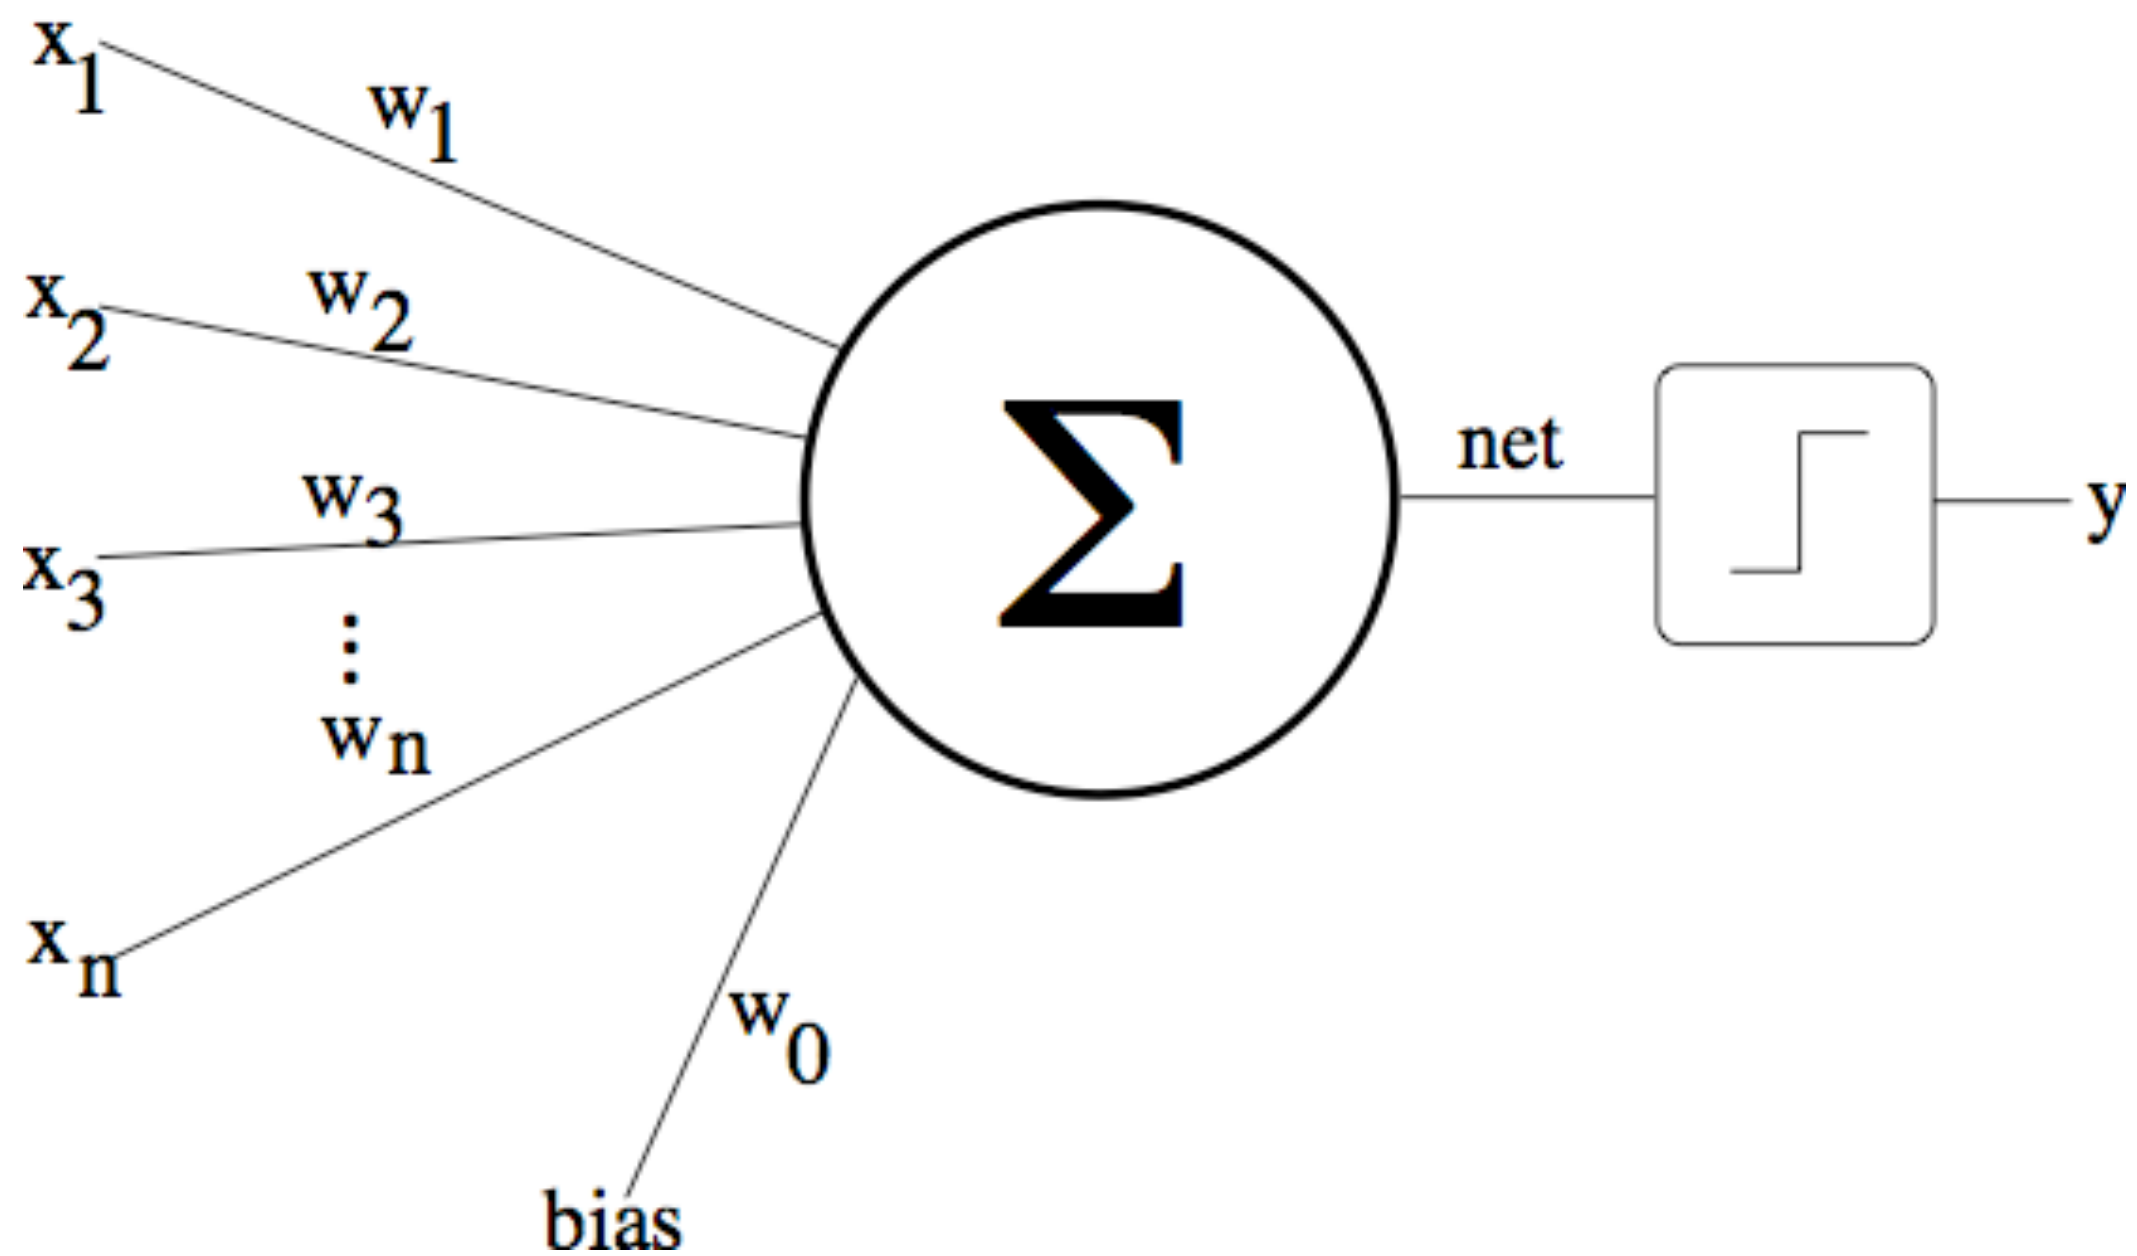
CS-GY-6613

Julian Togelius

[julian.togelius@nyu.edu](mailto:julian.togelius@nyu.edu)

# The Neuron









$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$

# Perceptron problems


- Noise: if the data isn't separable, weights might thrash
- Mediocre generalization: finds a “barely” separating solution
- Overtraining: test / held-out accuracy usually rises, then falls



# Which of these separators is optimal?




# Problem: learn this!




**With all of these problems and with so many other learning algorithms, why were people still interested in pursuing this technique?**


# Could we just train a bunch of perceptrons connected to each other?




$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$




$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$




$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$



$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$



$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$



$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$




# **Multilayer Perceptron**

**“When an axon of cell A is near enough cell B and repeatedly or persistently takes part in firing it, some growth process or metabolic change takes place in one or both cells such that A’s efficiency, as one of the cells firing B, is increased”**

*–Donald Hebb*

# Multi-layer Perceptron (MLP)



# Universal Approximator

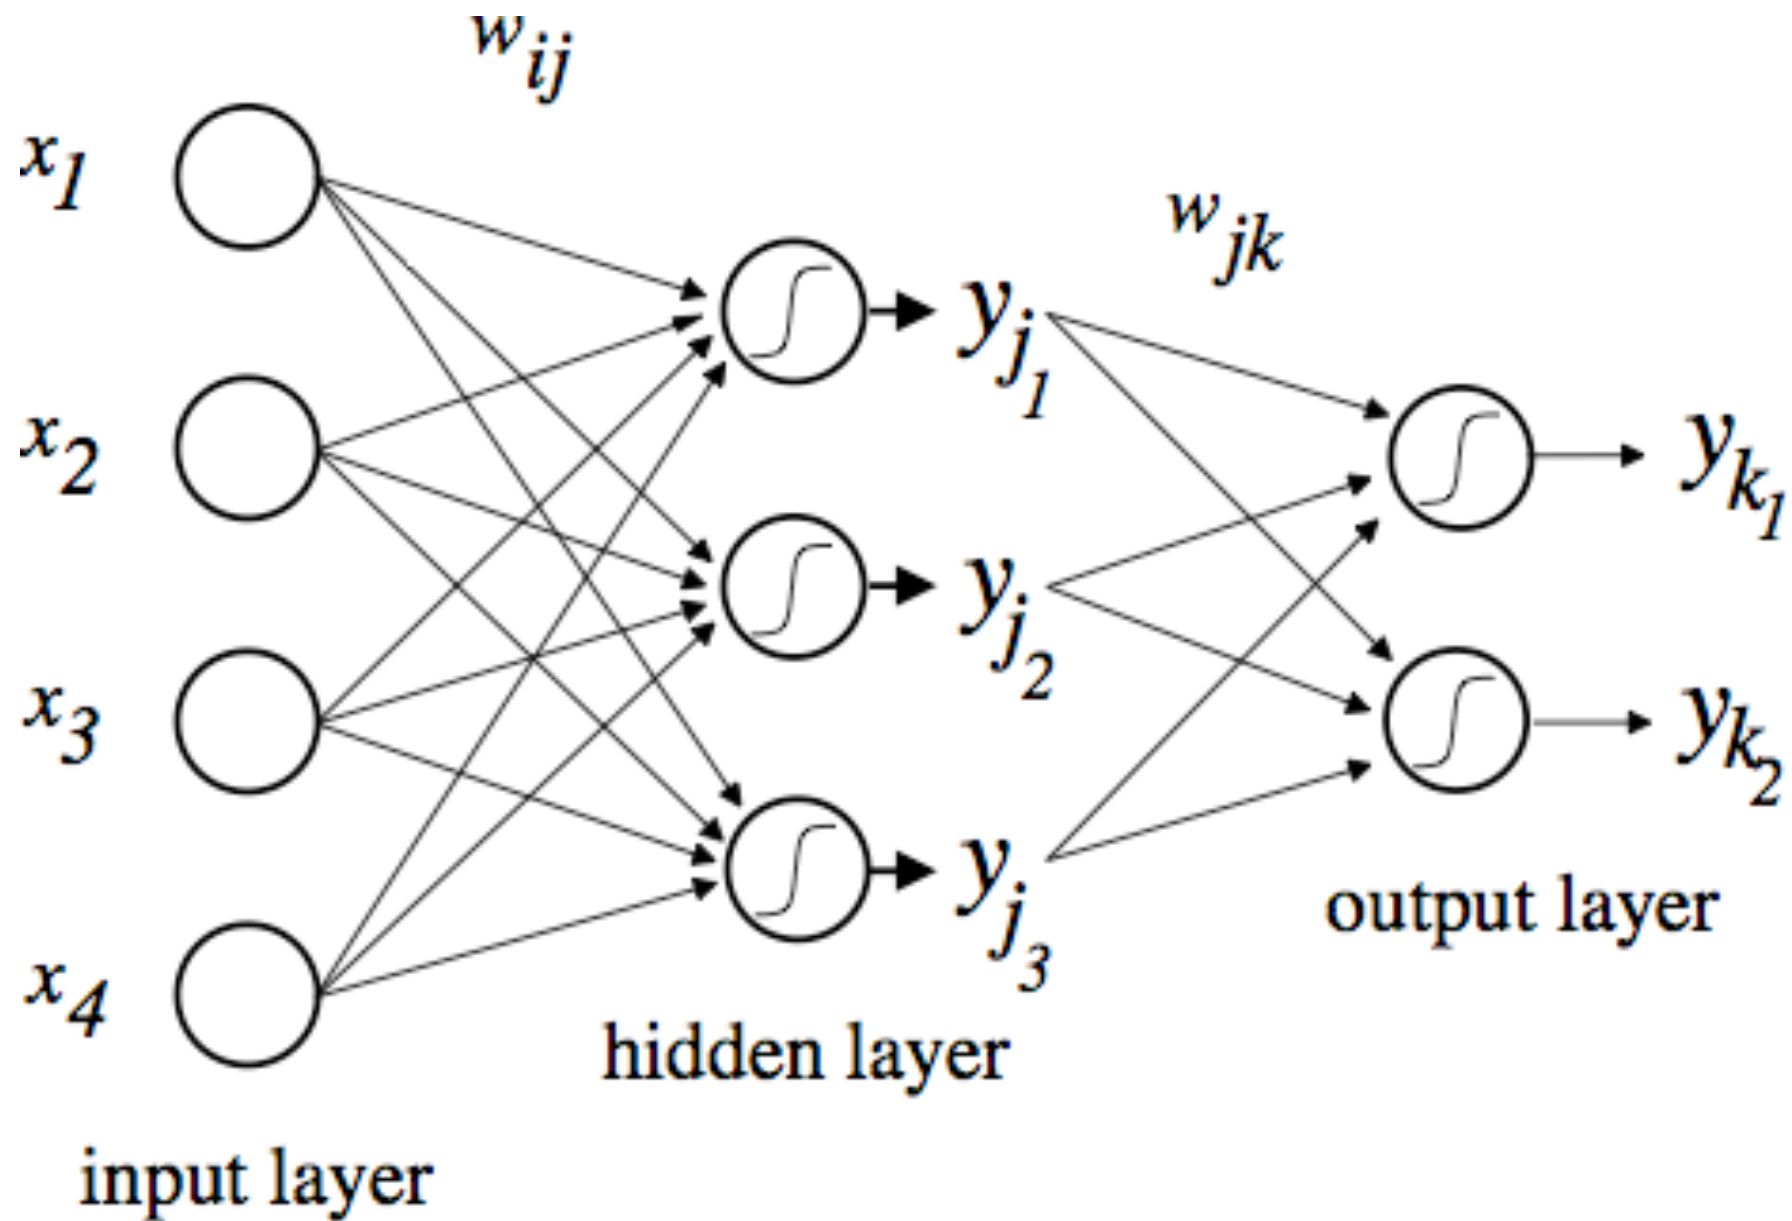
- George Cybenko in 1989 provided one of the first proofs for the Universal Approximation Theorem for Neural Networks
- States that a feedforward network with at least one hidden layer of a finite size can approximate any continuous function on compact subsets of  $\mathbb{R}^n$
- Says nothing about learnability
  - How do you train a hidden layer? What is its loss?

# Training a Neural Network


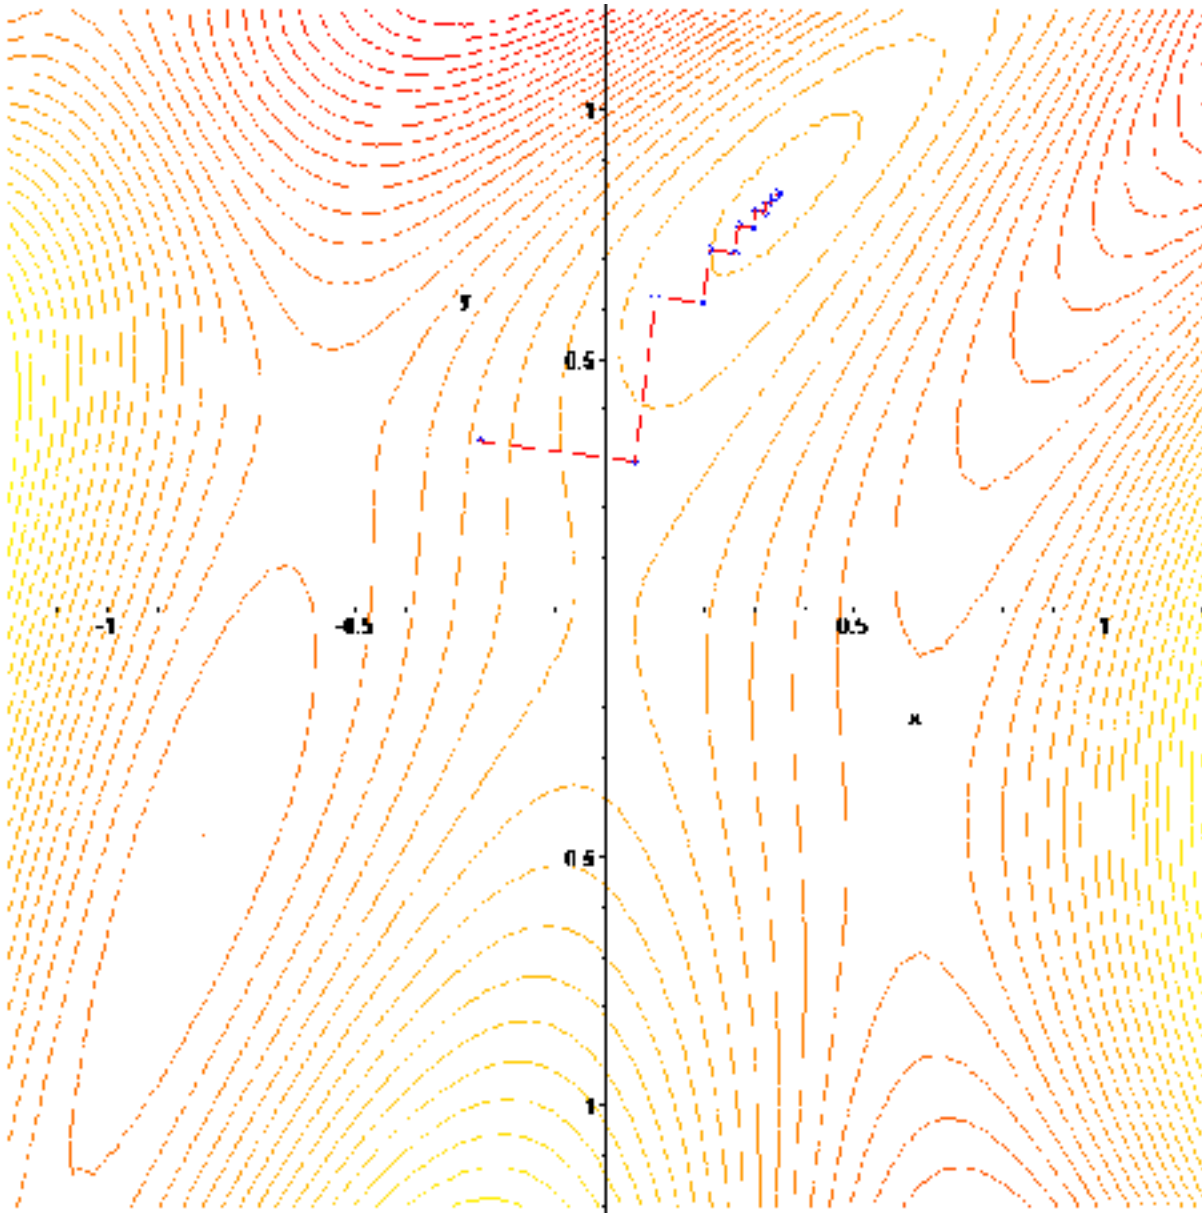
# Backpropagation

- Forward Pass: present training input pattern to network and activate network to produce output (can also do in batch: present all patterns in succession)
- Backward Pass: calculate error gradient and update weights starting at output layer and then going back

# Multi-layer Perceptron (MLP)



# Gradient Descent





# Gradient Descent

**Input** : list of  $n$  training examples  $(x_0, d_0) \dots (x_n, d_n)$

where  $\forall i : d_i \in \{+1, -1\}$

**Output** : classifying hyperplane  $w$

**Algorithm** :

Randomly initialize  $w$ ;

**While** makes errors on training set **do**

**for**  $(x_i, d_i)$  **do**

        let  $y_i = \text{MLP}(w, x_i)$ ;

$loss \leftarrow \text{Mean}((d_i - y_i)^2)$

$w' \leftarrow \text{Backprop}(w, loss)$

$w \leftarrow w - \eta w'$

**end**

**end**

*$x$  and  $w$  are vectors;  
 $i$  is the instance index*

# Building a Neural Network

## Weight Update Equation

$$\theta_w' = \theta_w - \eta * \frac{\partial TC}{\partial \theta_w}$$

**New weight**

**Learning Rate**

**function** BACK-PROP-LEARNING(*examples*, *network*) **returns** a neural network

**inputs:** *examples*, a set of examples, each with input vector  $\mathbf{x}$  and output vector  $\mathbf{y}$   
*network*, a multilayer network with  $L$  layers, weights  $w_{i,j}$ , activation function  $g$

**local variables:**  $\Delta$ , a vector of errors, indexed by network node

**repeat**


**for each** weight  $w_{i,j}$  in *network* **do**  
         $w_{i,j} \leftarrow$  a small random number

**for each** example  $(\mathbf{x}, \mathbf{y})$  in *examples* **do**  
        /\* Propagate the inputs forward to compute the outputs \*/  
        **for each** node  $i$  in the input layer **do**  
             $a_i \leftarrow x_i$   
        **for**  $\ell = 2$  to  $L$  **do**  
            **for each** node  $j$  in layer  $\ell$  **do**  
                 $in_j \leftarrow \sum_i w_{i,j} a_i$   
                 $a_j \leftarrow g(in_j)$   
        /\* Propagate deltas backward from output layer to input layer \*/  
        **for each** node  $j$  in the output layer **do**  
             $\Delta[j] \leftarrow g'(in_j) \times (y_j - a_j)$   
        **for**  $\ell = L - 1$  to  $1$  **do**  
            **for each** node  $i$  in layer  $\ell$  **do**  
                 $\Delta[i] \leftarrow g'(in_i) \sum_j w_{i,j} \Delta[j]$   
        /\* Update every weight in network using deltas \*/  
        **for each** weight  $w_{i,j}$  in *network* **do**  
             $w_{i,j} \leftarrow w_{i,j} + \alpha \times a_i \times \Delta[j]$

**until** some stopping criterion is satisfied

**return** *network*

# Activation Functions




$$f(x) = \sigma\left(\sum_{i=1}^n w_i x_i + b\right)$$

$$\sigma(net) = \frac{1}{1 + e^{-net}}$$

# Activation Functions


**Unit step (threshold)**

$$f(x) = \begin{cases} 0 & \text{if } 0 > x \\ 1 & \text{if } x \geq 0 \end{cases}$$




**Derivate: N/A**

# Activation Functions



**Derivative:  $\sigma'(z) = \sigma(z)(1 - \sigma(z))$**


# Activation Functions



**Hyperbolic Tangent**

**Derivative:  $\tanh'(z) = 1 - \tanh(z)^2$**

# Activation Functions

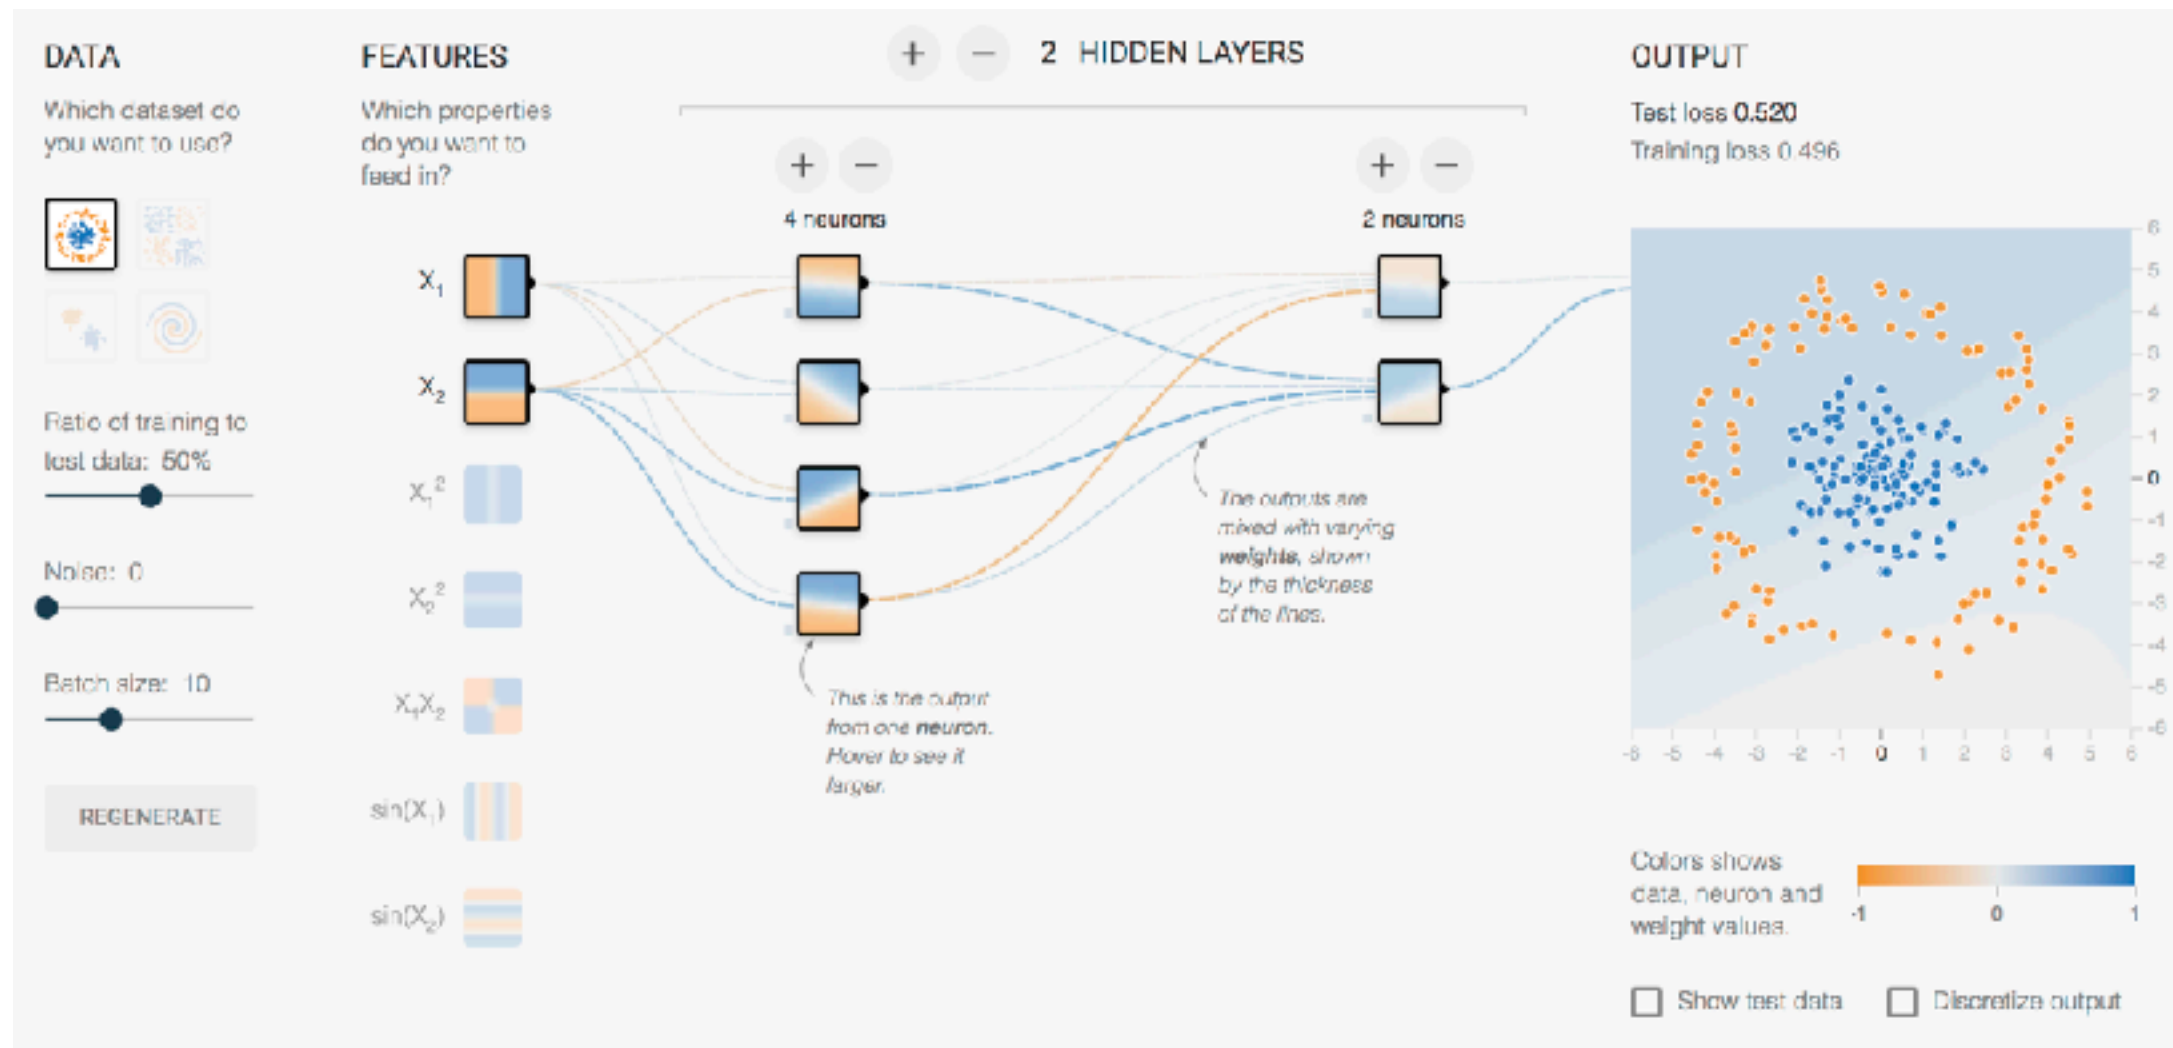


**Rectified Linear Units  
(ReLU)**

**Derivative: 1, if  $x > 0$ ; otherwise 0**



# Elements of a Neural Network



<http://playground.tensorflow.org/>

# Perceptron problems

- Noise: if the data isn't separable, weights might thrash
- Multiple layers allow it to work on non-linearly separable data
- Mediocre generalization: finds a “barely” separating solution
- More data and a continuous loss function greatly improve generalization
- Overtraining: test / held-out accuracy usually rises, then falls
- Large data, many layers, and batch training help

