

Lecture 23: Reinforcement Learning through Neuroevolution

Artificial Intelligence

Julian Togelius

Neuroevolution

- Training neural networks using evolutionary computation
- Representation: the weights of the neural network
- Fitness function: how well does the network perform its task?
- Works really well - when the network is small...


A tale from Ancient History

The challenge: car racing


- Driving a car fast requires fine motor control (in both senses)
- Optimizing lap times requires planning
- Overtaking requires adversarial planning




A simple car game

- Walls are solid
- Waypoints must be passed in order
- Fitness: continuous approximation of waypoints passed in 700 time steps



- Inputs
 - Six range-finder sensors (evolvable pos.)
 - Waypoint sensor, Speed, Bias
- Networks
 - Standard multi-layer perceptron, 9:6:2
 - Outputs interpreted as thrust/steering





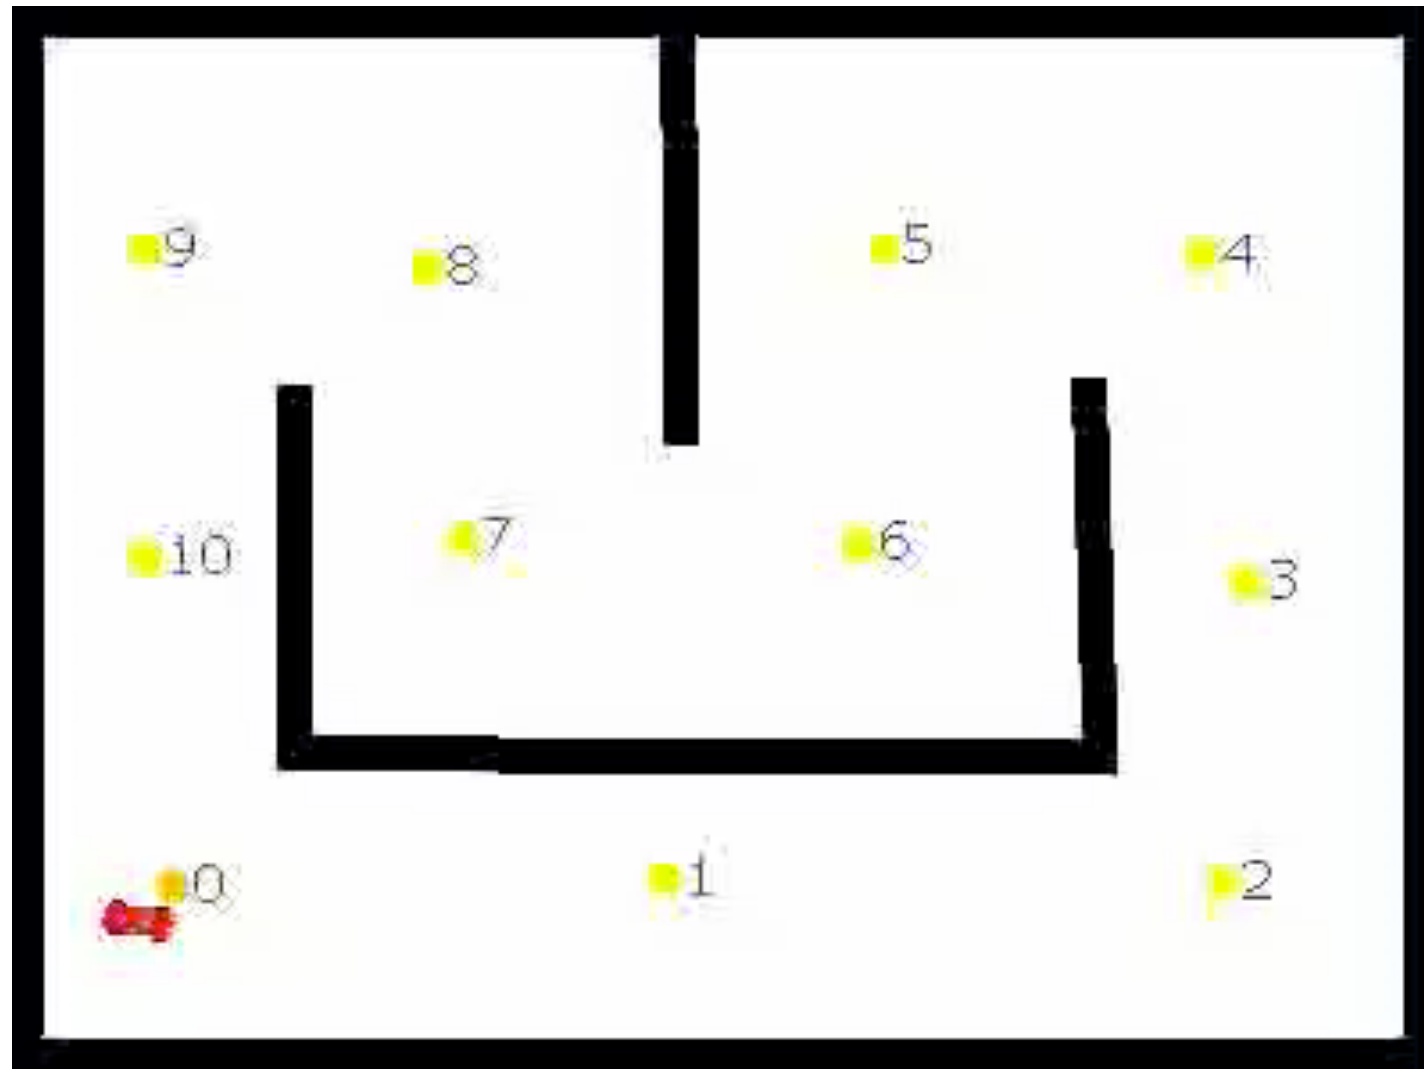
Algorithm 2: Evolution Strategy(μ, λ, n)

```
1 INITIALIZE (Population,  $\mu + \lambda$  individuals)
2 for  $i=1$  to  $n$  do
3   for  $j=1$  to  $(\mu + \lambda)$  do
4     EVALUATE (Population[j])
5   end
6   PERMUTE (Population)
7   SORTONFITNESS (Population)
8   for  $j=\mu$  to  $(\mu + \lambda)$  do
9     Population[j]  $\leftarrow$  COPY (Population[j- $\lambda$ ])
10    WEIGHTMUTATE (Population[j])
11  end
12 end
```

Mutation: add Gaussian noise with sd 1 to each connection

Fitness: progress around the track


Example video



Evolved with 50+50 ES, 100 Generations

Generalization and specialization


- A controller evolved for one track does not necessarily perform well on other tracks
 - In fact, a controller evolved to drive a track clockwise will not be able to drive *the same track* counterclockwise
- How do we achieve more general game-playing skills?
 - Is there a tradeoff between generality and performance?




Incremental evolution

- Introduced by Gomez & Mikkulainen (1997)
 - Later also called curriculum learning
- Change the fitness function f (to make it more demanding) as soon as a certain fitness is achieved
- In this case, add new tracks to f as soon as the controller can drive 1.5 rounds on all tracks currently in f

Incremental evolution



Video: navigating a complex track



Observations

- Controllers evolved for specific tracks perform poorly on other tracks
- General controllers, that can drive almost any track, can be incrementally evolved
- Starting from a general controller, a controller can be further evolved for *specialization* on a particular track
 - drive faster than the general controller
 - works even when evolution from scratch did not work!

Two cars on a track

- Two car with solo-evolved controllers on one track: disaster
 - they don't even see each other!
- How do we train controllers that take other drivers into account? (avoiding collisions or using them to their advantage)
- Solution: car sensors (rangefinders, like the wall sensors) and *competitive coevolution*


Competitive coevolution

- The fitness function evaluates at least two individuals
- One individual's success is *adversely* affected by the other's (directly or indirectly)
- Very potent, but seldom straightforward; e.g. Hillis (1991), Rosin and Belew (1996)


Competitive coevolution

- Standard 15+15 ES; each individual is evaluated through testing against the current best individual in the population
- Fitness function a mix of...
 - Absolute fitness: progress in n time steps
 - Relative fitness: distance ahead of or behind the other car after n time steps


Video: absolute fitness



Video: 50/50 fitness



Video: relative fitness



Is neuroevolution a good way of doing RL?

- In theory, it should be less efficient
- In practice, it is often more robust
- Plus, you can do multiobjective evolution, quality-diversity illumination, etc.