

# Lecture 17:

# Decision Trees

Artificial Intelligence

CS-GY-6613-I

Julian Togelius

[julian.togelius@nyu.edu](mailto:julian.togelius@nyu.edu)

# Types of learning

- **Supervised learning**

Learning to predict or classify labels based on labeled input data



- **Unsupervised learning**

Finding patterns in unlabeled data



- **Reinforcement learning**

Learning well-performing behavior from state observations and rewards

# Model construction



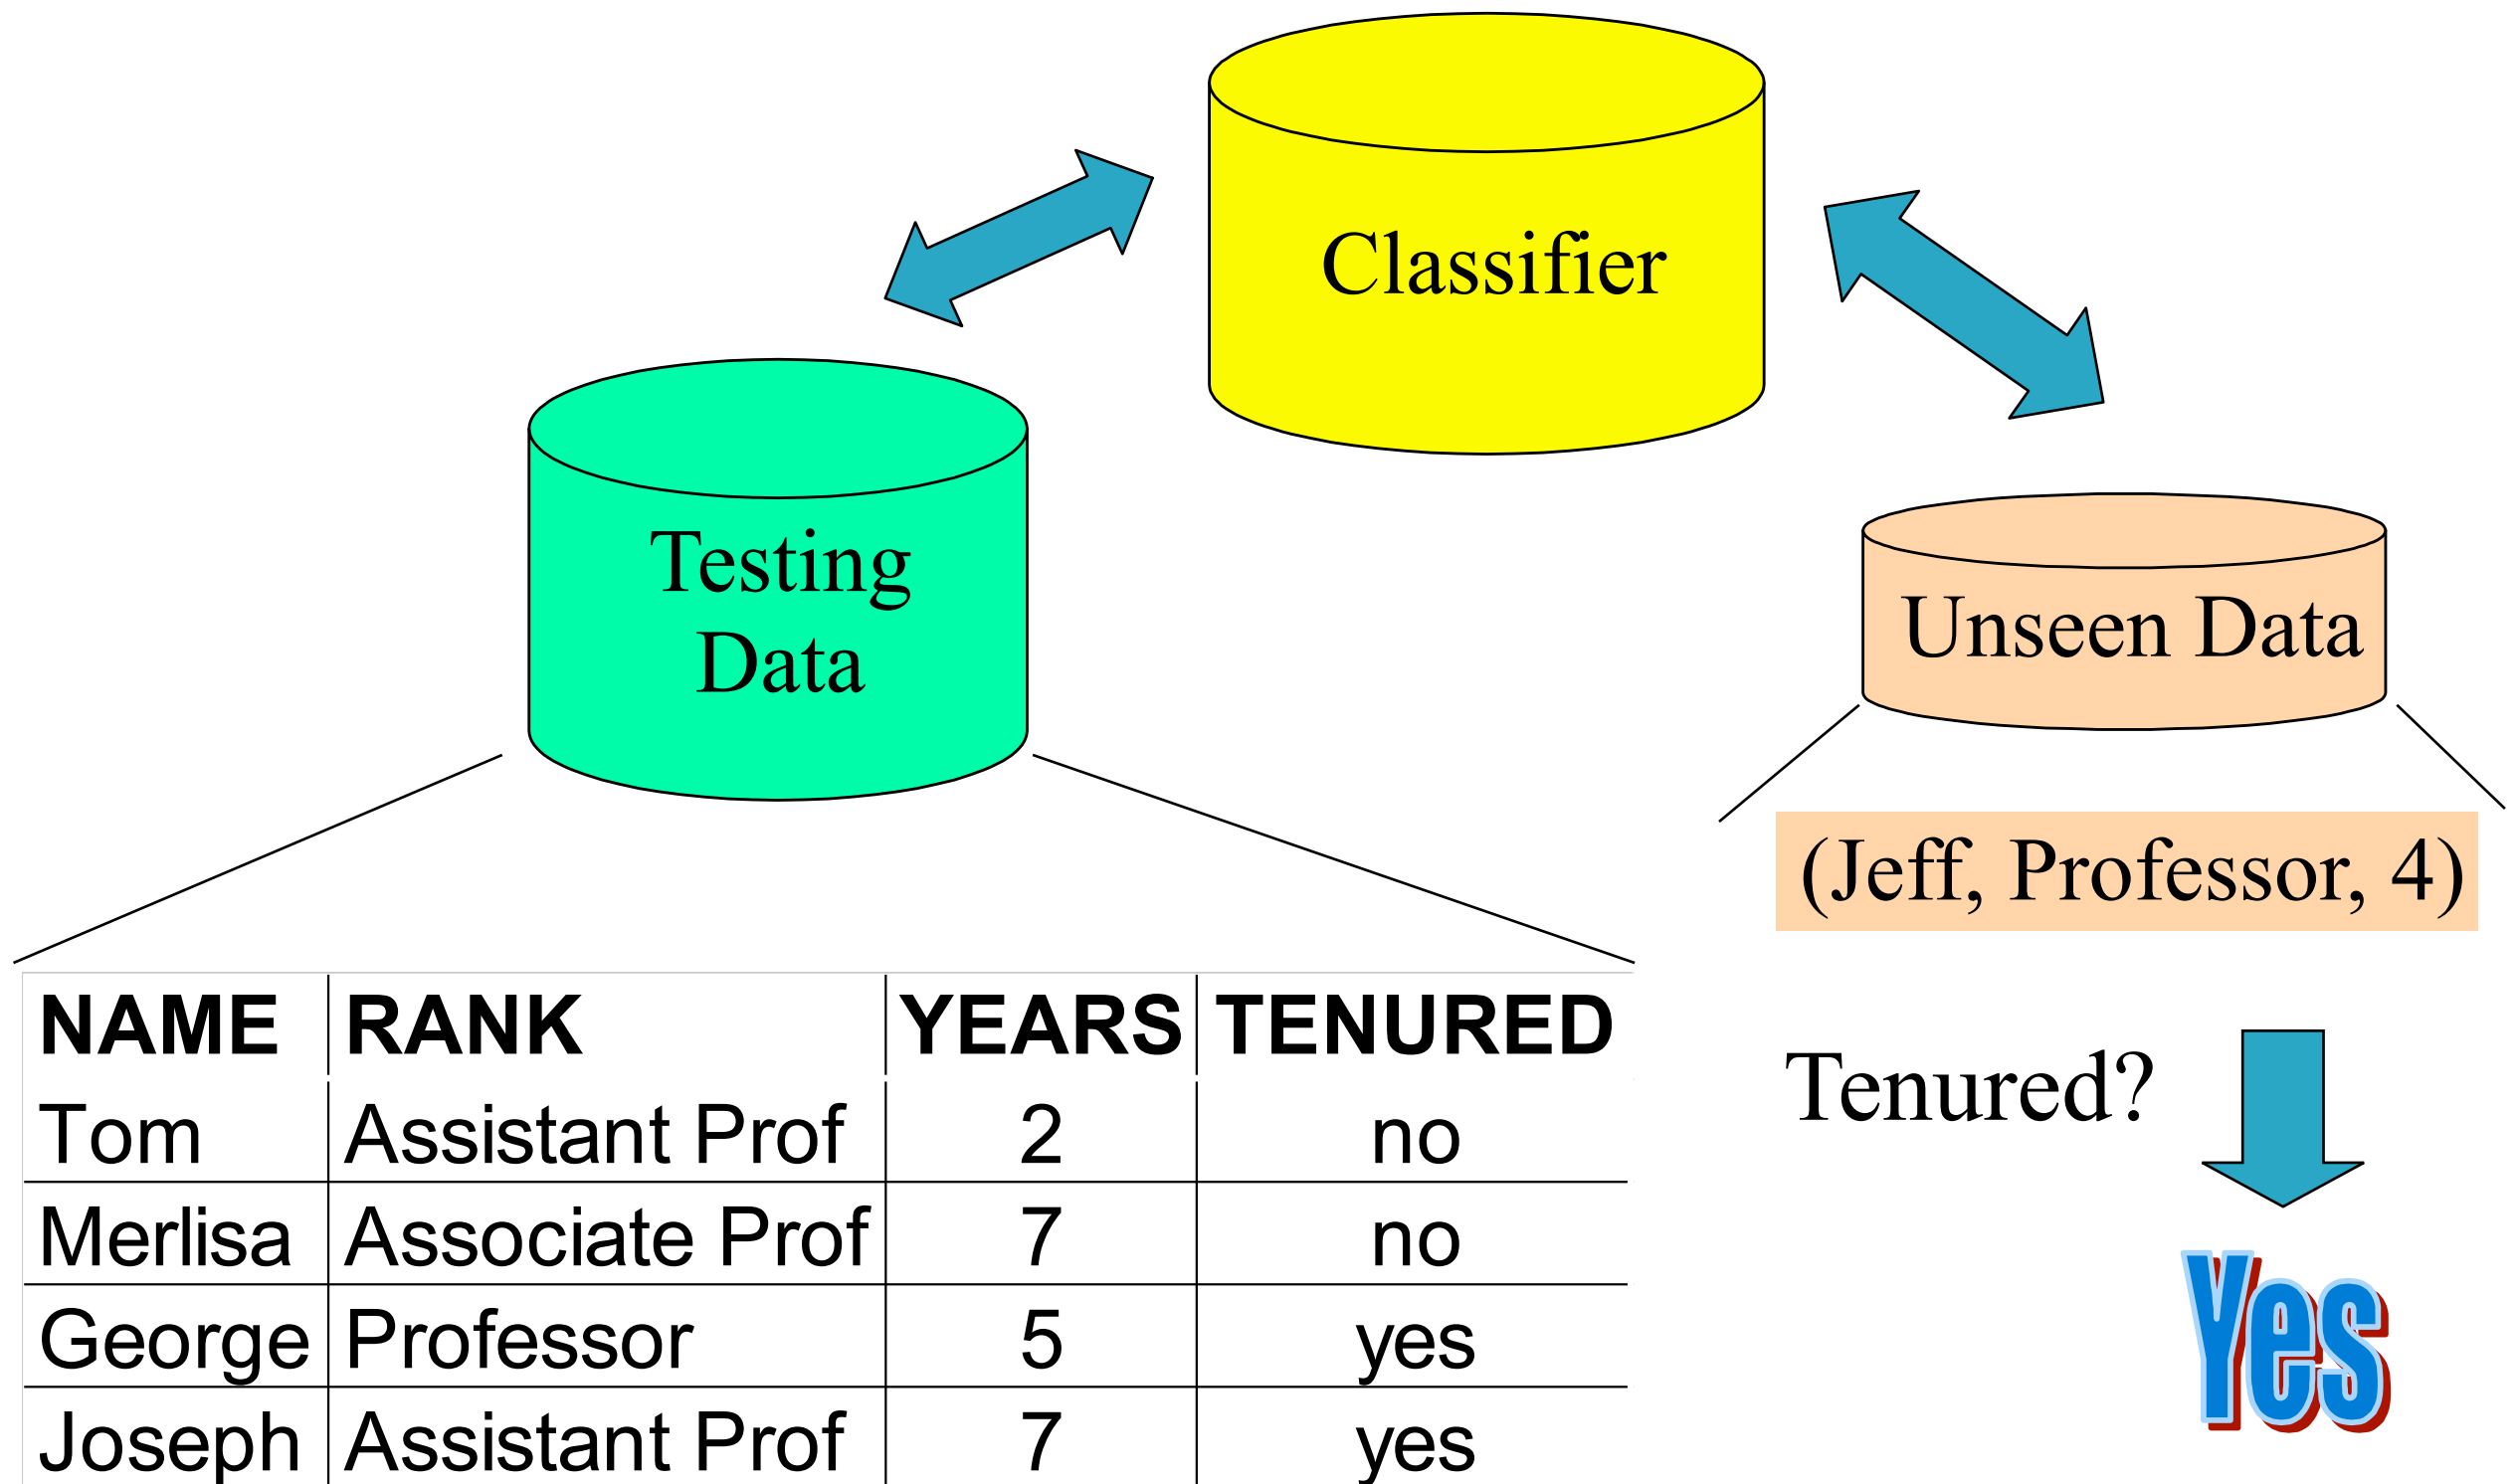
Classification  
Algorithms



NAME	RANK	YEARS	TENURED
Mike	Assistant Prof	3	no
Mary	Assistant Prof	7	yes
Bill	Professor	2	yes
Jim	Associate Prof	7	yes
Dave	Assistant Prof	6	no
Anne	Associate Prof	3	no

IF rank = 'professor'  
OR years > 6  
THEN tenured = 'yes'

# Using the model



# Classification vs prediction

- Classification: binary or nominal labels
  - Examples: pregnant or not, from which country, which type of road sign
- Prediction: continuous labels
  - Examples: future stock price, life expectancy, distance to obstacle

# Terminology (supervised learning)

- Each line of data: instance / data point / tuple
- The features of each instance: features / attributes
- That which should be learned: labels / targets
- Each instance has features and a label
- We train on the training set...
- ...and test on the testing set

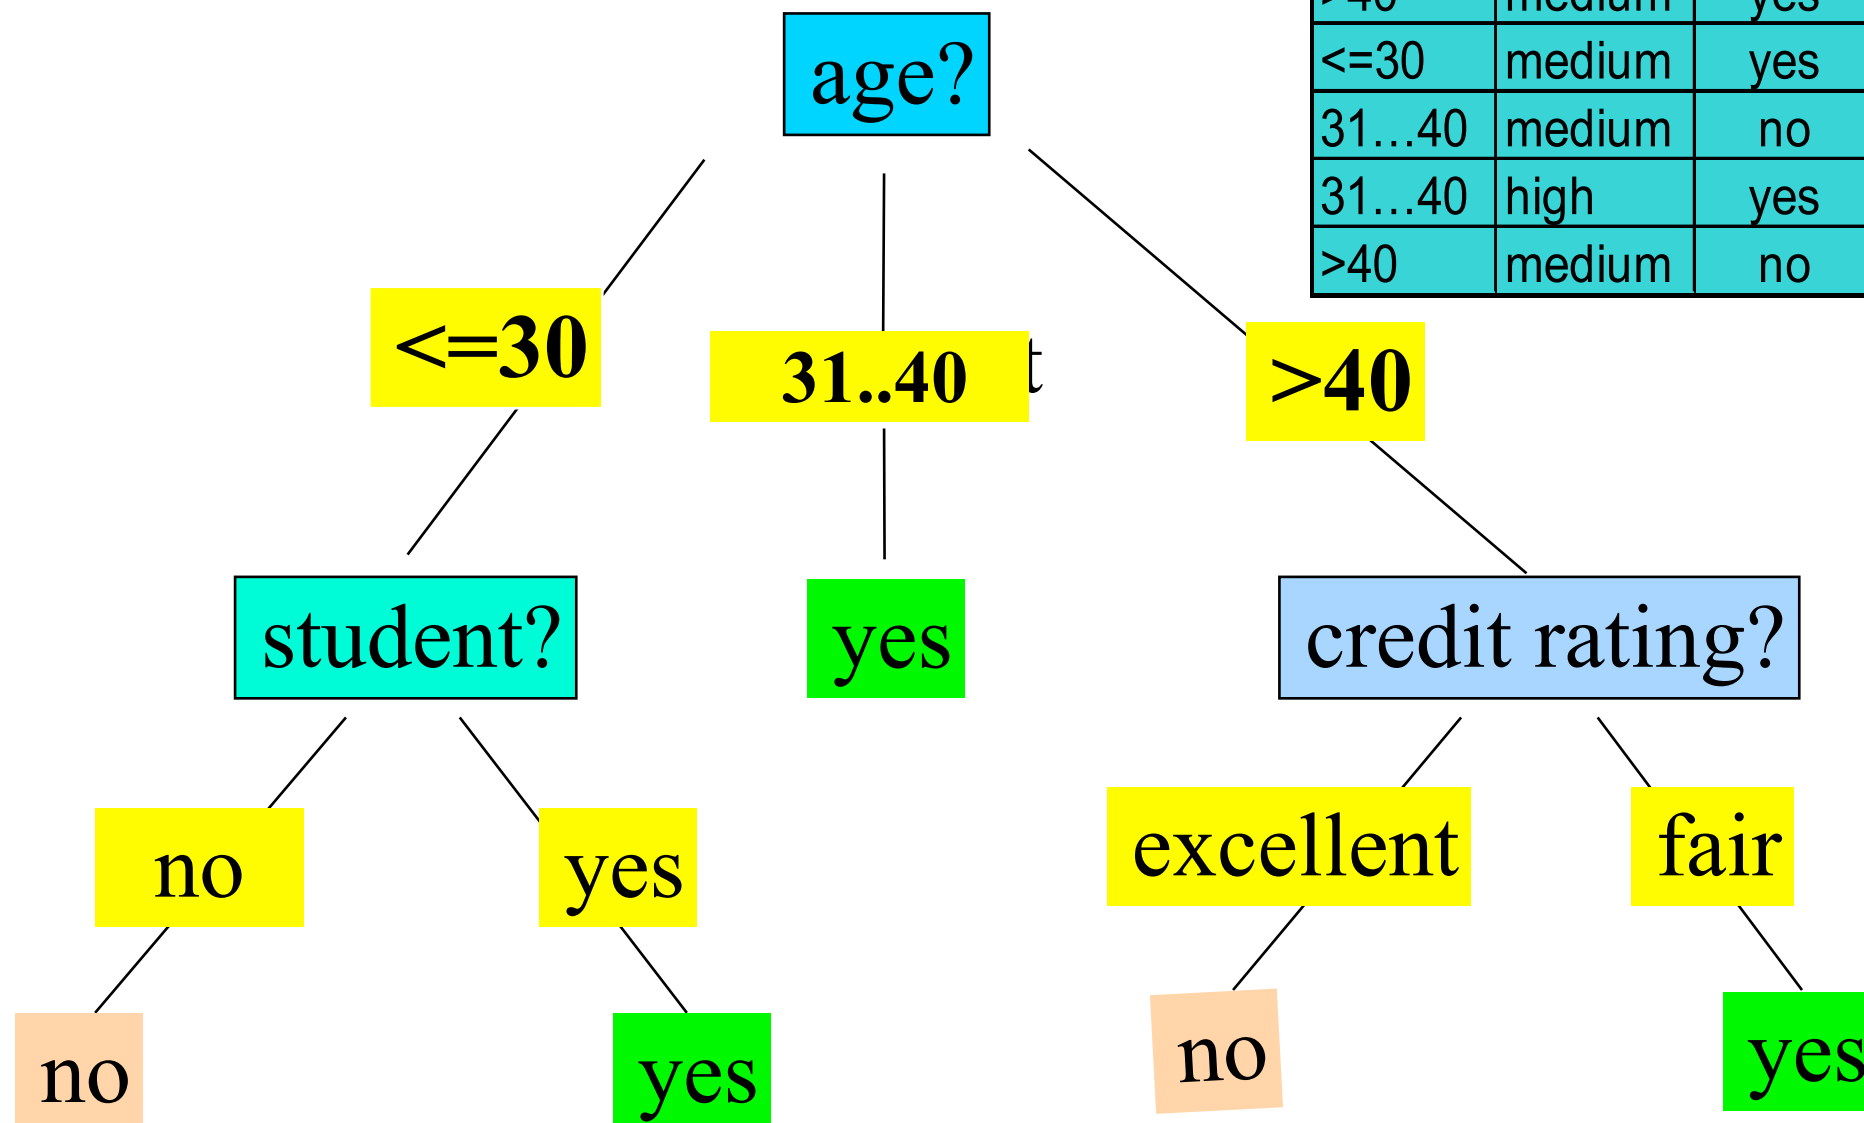
# Decision trees

- A popular representation for classifiers
- Human-readable
- Can be learned (induced) efficiently using algorithms based on information theory
  - An eager learning method
- Often yields high-accuracy classifiers

# An example

Classify: buys\_computer

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no





# ID3 Algorithm for Decision Tree Induction

- Tree is constructed in a top-down recursive divide-and-conquer manner
- At start, all the training examples are at the root
- Attributes are categorical (if continuous-valued, they are discretized in advance)
- Examples are partitioned recursively based on selected attributes
- Test attributes are selected on the basis of a heuristic or statistical measure (e.g., information gain)

# ID3 Algorithm for Decision Tree Induction

- Stop partitioning when...
  - All samples for a given node belong to the same class, or...
  - There are no remaining attributes for further partitioning – (vote on the leaf) or...
  - There are no samples left

# Attribute Selection Measure: Information Gain (ID3/C4.5)

- Select the attribute with the highest information gain
- Let  $p_i$  be the probability that an arbitrary tuple in  $D$  belongs to class  $C_i$ , estimated by  $|C_i, D|/|D|$

- Expected information (entropy) needed to classify a tuple in  $D$ :

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$


- Information needed (after using  $A$  to split  $D$  into  $v$  partitions) to classify  $D$ :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

- Information gained by branching on attribute  $A$

$$Gain(A) = Info(D) - Info_A(D)$$

# Binary entropy function



# Attribute Selection: Information Gain

Class P: buys\_computer = “yes”

Class N: buys\_computer = “no”

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

age	$p_i$	$n_i$	$I(p_i, n_i)$
$\leq 30$	2	3	0.971
31...40	4	0	0
$> 40$	3	2	0.971

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

$\frac{5}{14} I(2,3)$  means “age  $\leq 30$ ” has 5 out of 14 samples, with 2 yes’es and 3 no’s. Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly,

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit\_rating) = 0.048$$

age	income	student	credit_rating	buys_computer
$\leq 30$	high	no	fair	no
$\leq 30$	high	no	excellent	no
31...40	high	no	fair	yes
$> 40$	medium	no	fair	yes
$> 40$	low	yes	fair	yes
$> 40$	low	yes	excellent	no
31...40	low	yes	excellent	yes
$\leq 30$	medium	no	fair	no
$\leq 30$	low	yes	fair	yes
$> 40$	medium	yes	fair	yes
$\leq 30$	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
$> 40$	medium	no	excellent	no

- Input: A data set,  $S$   
Output: A decision tree
  - If all the instances have the same value for the target attribute then return a decision tree that is simply this value (not really a tree - more of a stump).
- Else
  - Compute Gain values (see above) for all attributes and select an attribute with the highest value and create a node for that attribute.
  - Make a branch from this node for every value of the attribute
  - Assign all possible values of the attribute to branches.
  - Follow each branch by partitioning the dataset to be only instances whereby the value of the branch is present and then go back to 1.

**function** DECISION-TREE-LEARNING(*examples*, *attributes*, *parent\_examples*) **returns**  
tree

**if** *examples* is empty **then return** PLURALITY-VALUE(*parent\_examples*)  
**else if** all *examples* have the same classification **then return** the classification  
**else if** *attributes* is empty **then return** PLURALITY-VALUE(*examples*)  
**else**

$A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \text{IMPORTANCE}(a, \text{examples})$

*tree*  $\leftarrow$  a new decision tree with root test *A*

**for each** value  $v_k$  of *A* **do**


$\text{exs} \leftarrow \{e : e \in \text{examples} \text{ and } e.A = v_k\}$

*subtree*  $\leftarrow$  DECISION-TREE-LEARNING(*exs*, *attributes* – *A*, *examples*)


add a branch to *tree* with label (*A* =  $v_k$ ) and subtree *subtree*

**return** *tree*


# Overfitting (in supervised learning)



Underfitted



Good Fit/Robust



Overfitted

- Fitting the data too well, including the noise
- Reduces accuracy on unseen data



# Overfitting and pruning

- Overfitting: An induced tree may overfit the training data (specific instance of a concept that applies to all supervised learning)
  - Too many branches, some may reflect anomalies due to noise or outliers
  - Poor accuracy for unseen samples
- Three approaches to avoid overfitting
  - Prepruning: Halt tree construction early— do not split a node if this would result in the goodness measure (e.g. accuracy on a validation set) falling below a threshold
  - Postpruning: Remove branches from a “fully grown” tree—get a sequence of progressively pruned trees
    - Use a set of data different from the training data to decide which is the “best pruned tree”
  - Limit depth or size

# What to do with numerical attributes?

- Create a set of bins, where each bin becomes one attribute value (e.g. income <30k, 30-40k, >40k)
- Equal depth: All the bins have the same number of instances in them
- Equal width: All the bins have the same range
- More sophisticated ways, including looking at where there are big discontinuities in the range of the input data etc

# Can we do better?

- It's hard to train good classifiers
- Can we replace quality with quantity?
  - Train many classifiers and combine them somehow
  - Called *ensemble learning*